

▼ 1.1) Without Data NOrmalization / Standardization

```
from google.colab import drive
drive.mount('/content/drive')

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount)

import pandas as pd

data = pd.read_csv('/content/drive/My Drive/ISL/advertising.csv')

data.head()

TV  Radio  Newspaper  Sales
0  230.1    37.8      69.2   22.1
1   44.5    39.3      45.1   10.4
2   17.2    45.9      69.3   12.0
3  151.5    41.3      58.5   16.5
4  180.8    10.8      58.4   17.9

features = data[['TV', 'Radio', 'Newspaper']].values
```

▼ 3-D tsn3 visualization

```
from sklearn.manifold import TSNE

tsne = TSNE(n_components=3, perplexity=30, n_iter=3000)
reduced_features = tsne.fit_transform(features)

import matplotlib.pyplot as plt
from mpl_toolkits.mplot3d import Axes3D

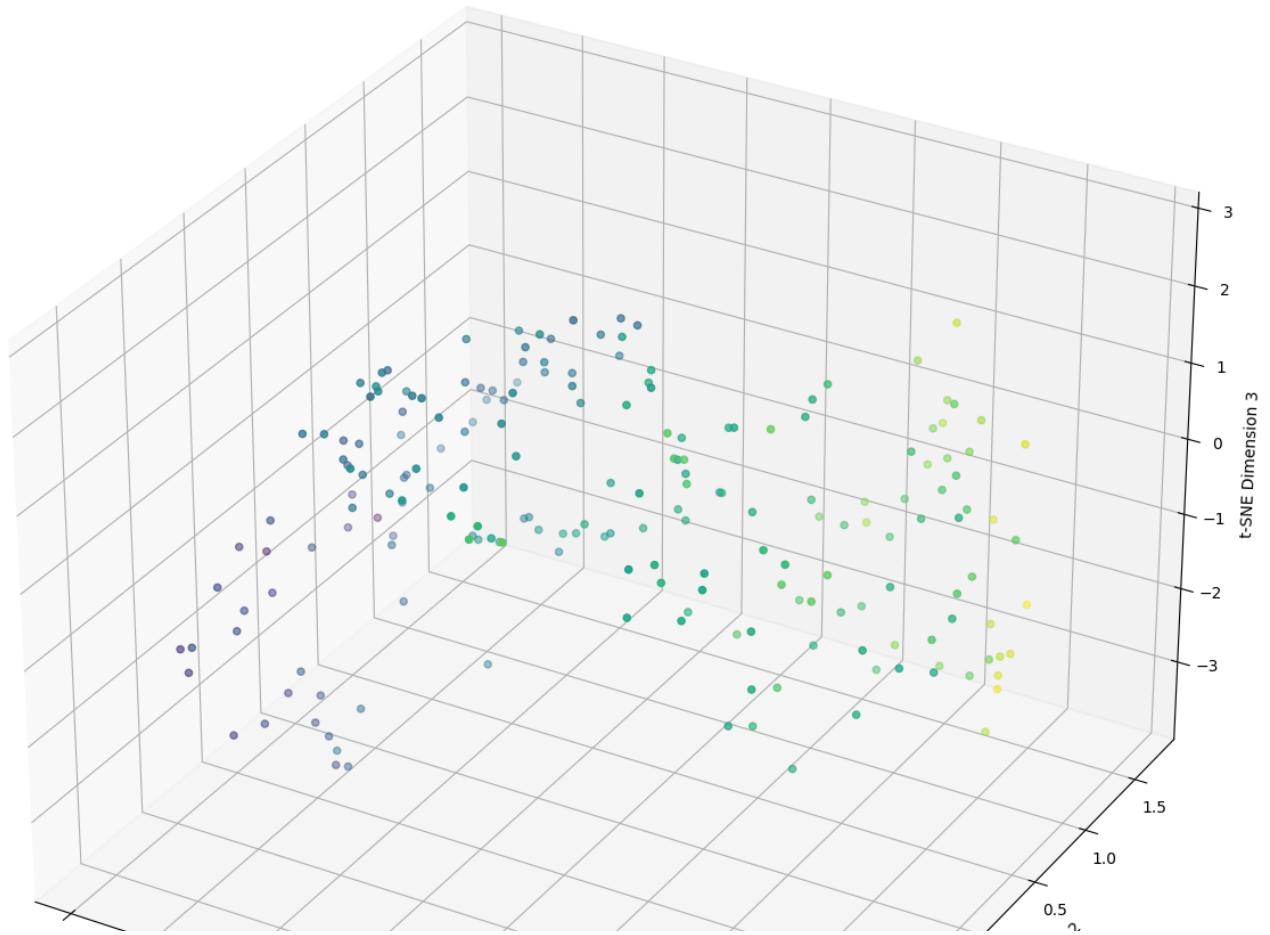
fig = plt.figure(figsize=(15,20))
ax = fig.add_subplot(111, projection='3d')

sales = data['Sales'].values

ax.scatter(reduced_features[:, 0], reduced_features[:, 1], reduced_features[:, 2], c=sales, cmap='viridis')
ax.set_xlabel('t-SNE Dimension 1')
ax.set_ylabel('t-SNE Dimension 2')
ax.set_zlabel('t-SNE Dimension 3')
ax.set_title('t-SNE Visualization of Data')

plt.show()
```

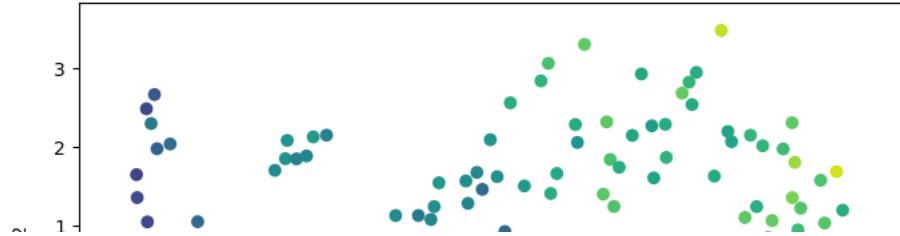
t-SNE Visualization of Data



▼ 2-D tsne visualtion

```
-> 5      X      / \      / /      / / -0.5 ↗  
tsne = TSNE(n_components=2, perplexity=30, n_iter=300) # for 2D visualization  
  
reduced_features = tsne.fit_transform(features)  
  
plt.figure(figsize=(8, 6))  
  
sales = data['Sales'].values  
  
plt.scatter(reduced_features[:, 0], reduced_features[:, 1], c=sales, cmap='viridis')  
plt.xlabel('t-SNE Dimension 1')  
plt.ylabel('t-SNE Dimension 2')  
plt.title('t-SNE 2D Visualization of Data')  
  
plt.show()
```

t-SNE 2D Visualization of Data



▼ Modelling

```
# Splitting the data into training and test data set , for this assignment taking the first 80 % samples as training and las
total = len(features)

train_no = int( 0.8 * total )
test_no = total - train_no

X = features[0:train_no]
y = data['Sales'].values[0:train_no]

test_features = features[train_no : 201 ]
test_labels = data['Sales'].values[train_no : 201]
```

X.shape

(160, 3)

```
import numpy as np
# Padding of 1 to matrix X
ones = np.ones((X.shape[0], 1))
X_padded = np.hstack((X, ones))
```

▼ Linear Model without Regularization

$$\hat{\beta}_{ls} = (X'X)^{-1}X'Y$$

import numpy as np

```
# As we know the parameter for linear regression if supposed as B then B = (X.T X) inv Xt y
X_padded = np.array(X_padded)
y = np.array(y)

B_linear = np.linalg.inv( X_padded.T @ X_padded ) @ ( X_padded.T @ y )

array([5.47207688e-02, 1.04776019e-01, 3.06638052e-03, 4.48194957e+00])
```

```
from sklearn.linear_model import LinearRegression
model = LinearRegression()
model.fit(X, y)
```

```
# Get the coefficients (parameters)
intercept = model.intercept_
coefficients = model.coef_

print("Intercept:", intercept)
print("Coefficients:", coefficients)
```

```
Intercept: 4.48194956663138
Coefficients: [0.05472077 0.10477602 0.00306638]
```

As we can see the coefficients values are exactly same as the one calculated by inbuilt libraries

X_padded.shape

(160, 4)

▼ linear Model with regularizartion (Ridge Regression)

$$\hat{\beta} = \hat{\beta}_{ridge} = (X'X + \lambda I_p)^{-1} X'Y$$

```
# As we know for Ridge regression the equation looks like the above
lemda = 0.01
B_ridge = np.linalg.inv( X_padded.T @ X_padded + lemda * np.identity(4) ) @ ( X_padded.T @ y )
B_ridge
```

```
array([5.47259510e-02, 1.04799327e-01, 3.07865631e-03, 4.47996558e+00])
```

```
from sklearn.linear_model import Ridge
alpha = 0.01 # You can adjust the alpha value
ridge_model = Ridge(alpha=alpha)
ridge_model.fit(X, y)

# Get the coefficients (parameters)
coefficients = ridge_model.coef_
intercept = ridge_model.intercept_

print("Intercept:", intercept)
print("Coefficients:", coefficients)
```

```
Intercept: 4.4819501842851786
Coefficients: [0.05472077 0.10477598 0.00306639]
```

As we can see the parameters are same for

The one calculated statistically

calculated with inbuilt libraires

▼ Computing errors

```
print(test_features.shape)
test_labels.shape
```

```
(40, 3)
(40,)
```

```
test_features = np.array(test_features)
test_labels = np.array(test_labels)

# padding test features with 1

ones = np.ones((test_features.shape[0] , 1))
test_features_padded = np.hstack((test_features , ones))
```

```
linear_pred_train = X_padded @ B_linear
ridge_pred_train = X_padded @ B_ridge
```

```
linear_pred_test = test_features_padded @ B_linear
ridge_pred_test = test_features_padded @ B_ridge
```

```
linear_train_error_1 = ((y - linear_pred_train).T @ ( y - linear_pred_train )) / ( 2 * ( X_padded.shape[0])) 
ridge_train_error_1 = ((y - ridge_pred_train).T @ ( y - ridge_pred_train )) / ( 2 * ( X_padded.shape[0]))
```

```
linear_test_error_1 = ((test_labels - linear_pred_test).T @ ( test_labels - linear_pred_test )) / ( 2 * ( test_features_padded
ridge_test_error_1 = ((test_labels - ridge_pred_test).T @ ( test_labels - ridge_pred_test )) / ( 2 * ( test_features_padded
```

```
print("Linear Regression Training Error:", linear_train_error_1)
```

```

print("Ridge Regression Training Error:", ridge_train_error_1)
print("Linear Regression Testing Error:", linear_test_error_1)
print("Ridge Regression Testing Error:", ridge_test_error_1)

Linear Regression Training Error: 1.3671985804792106
Ridge Regression Training Error: 1.3671988581486008
Linear Regression Testing Error: 1.3184703580408181
Ridge Regression Testing Error: 1.3186815950080715

```

▼ Experimenting with different values of lemda

```

lemdas = [ 0.01, 0.02, 0.03, 0.04, 0.05, 0.06, 0.07, 0.08, 0.09, 0.10, 0.11, 0.12, 0.13, 0.14, 0.15, 0.16, 0.17, 0.18, 0.19,
test_errors = []
train_errors = []
for lemda in lemdas:
    B_ridge = np.linalg.inv( X_padded.T @ X_padded + lemda * np.identity(4) ) @ ( X_padded.T @ y )

    pred = test_features_padded @ B_ridge
    test_error = ((( test_labels - pred).T)@ (test_labels - pred) ) / ( 2 * test_features_padded.shape[0])
    test_errors.append(test_error)

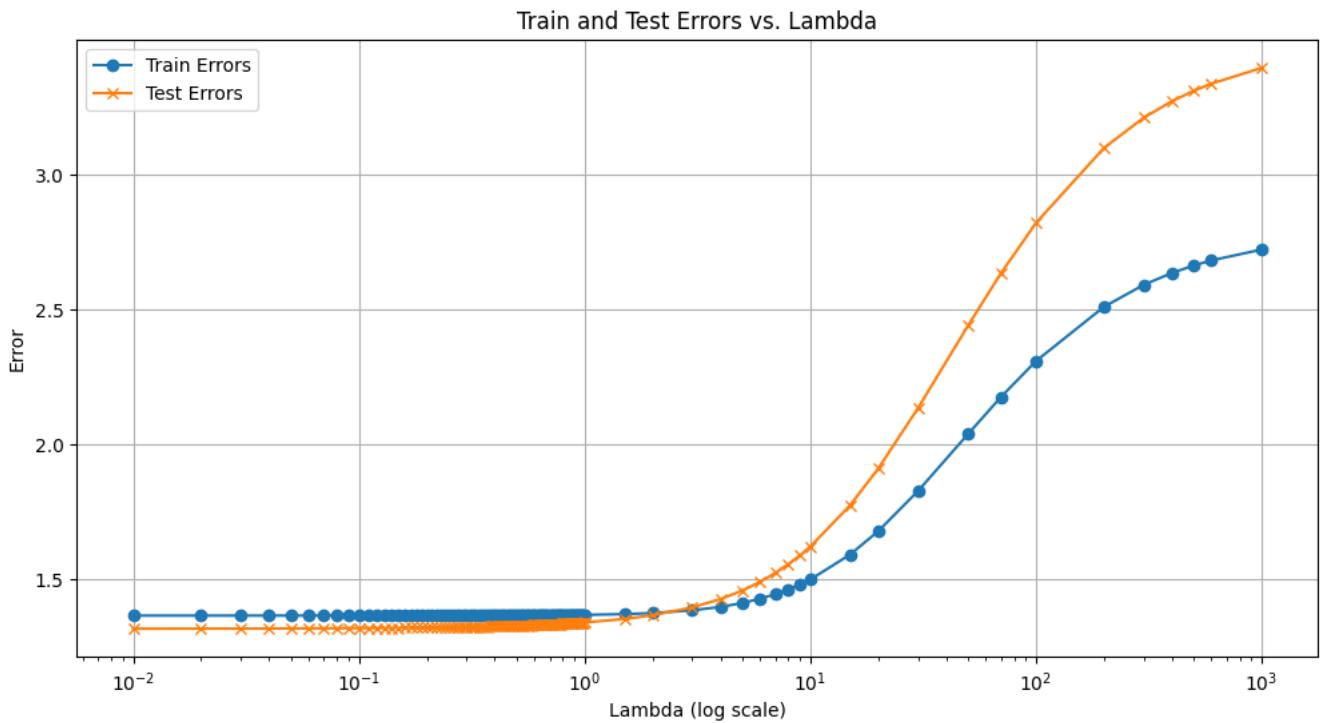
    pred = X_padded @ B_ridge
    train_error = ((( y - pred).T) @ ( y - pred) ) / ( 2 * X_padded.shape[0])
    train_errors.append(train_error)

```

```

# Create the plot
plt.figure(figsize=(12, 6))
plt.plot((lemdas), train_errors, label='Train Errors', marker='o')
plt.plot((lemdas), test_errors, label='Test Errors', marker='x')
plt.xscale('log') # To make the x-axis logarithmic
plt.xlabel('Lambda (log scale)')
plt.ylabel('Error')
plt.title('Train and Test Errors vs. Lambda')
plt.legend()
plt.grid(True)
plt.show()

```



#

▼ 1.2) With Normalization and standardizaion

```
min_vals = np.min(X, axis=0)
max_vals = np.max(X, axis=0)

min_vals_y = np.min(y , axis = 0 )
max_vals_y= np.max(y , axis=0)

X = (X - min_vals) / (max_vals - min_vals)
y = (y - min_vals_y) / (max_vals_y - min_vals_y)

mean_vals = np.mean(X, axis=0)
std_devs = np.std(X, axis=0)

mean_vals_y = np.mean(y, axis=0)
std_devs_y = np.std(y, axis=0)

X = (X - mean_vals) / std_devs
y = ( y - mean_vals_y) / std_devs_y
```

▼ 3-D tsne

```
from sklearn.manifold import TSNE

tsne = TSNE(n_components=3, perplexity=30, n_iter=3000)
reduced_features = tsne.fit_transform(features)

import matplotlib.pyplot as plt
from mpl_toolkits.mplot3d import Axes3D

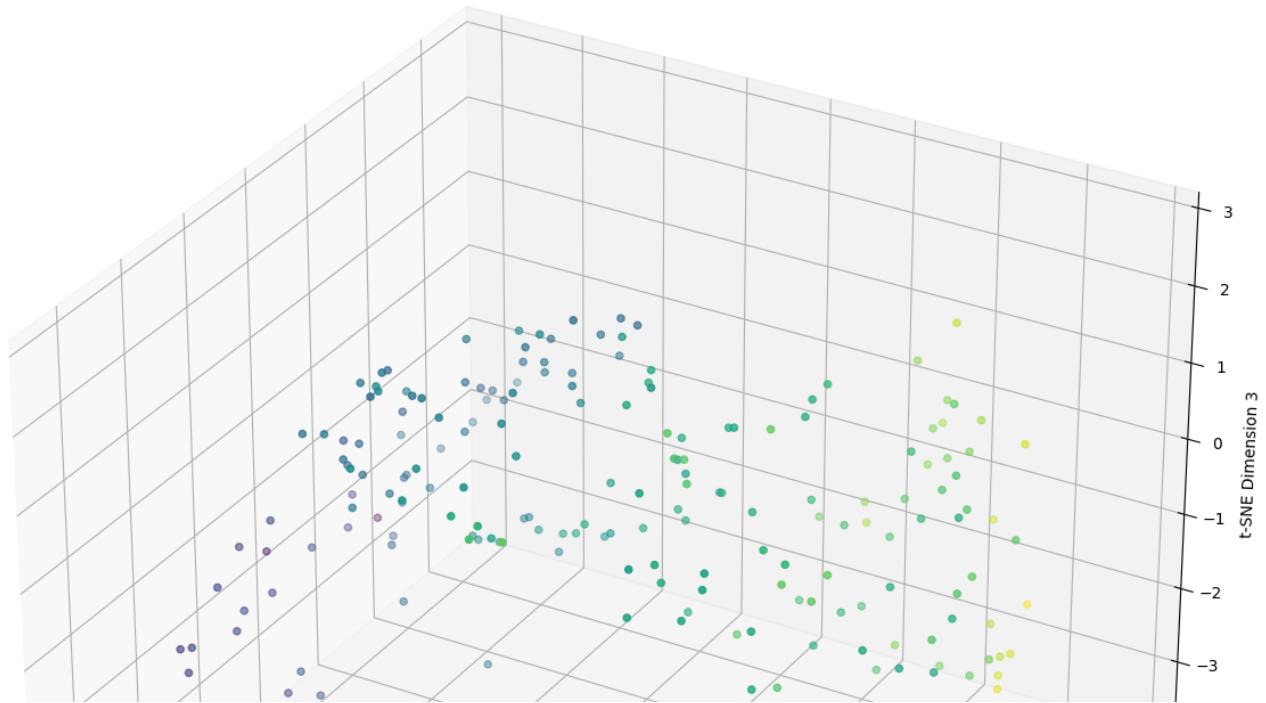
fig = plt.figure(figsize=(15,20))
ax = fig.add_subplot(111, projection='3d')

sales = data['Sales'].values

ax.scatter(reduced_features[:, 0], reduced_features[:, 1], reduced_features[:, 2], c=sales, cmap='viridis')
ax.set_xlabel('t-SNE Dimension 1')
ax.set_ylabel('t-SNE Dimension 2')
ax.set_zlabel('t-SNE Dimension 3')
ax.set_title('t-SNE Visualization of Data')

plt.show()
```

t-SNE Visualization of Data



▼ Padding of X

```
# Padding of 1 to matrix X
ones = np.ones((X.shape[0], 1))
X_padded = np.hstack((X, ones))
```

▼ Linear

```
X_padded = np.array(X_padded)
y = np.array(y)

B_linear = np.linalg.inv( X_padded.T @ X_padded ) @ ( X_padded.T @ y )
B_linear
```

```
array([8.86190469e-01, 2.90834750e-01, 1.27825629e-02, 3.88865703e-16])
```

▼ Ridge

```
lemda = 0.01
B_ridge = np.linalg.inv( X_padded.T @ X_padded + (lemda * np.identity(4)) ) @ ( X_padded.T @ y )

B_ridge
```

```
array([8.86135872e-01, 2.90816822e-01, 1.27899138e-02, 3.88832274e-16])
```

▼ Comparison of errors

```
test_features = (test_features - min_vals) / (max_vals - min_vals)
test_features = (test_features - mean_vals) / std_devs

test_labels = (test_labels - min_vals_y) / (max_vals_y - min_vals_y)
test_labels = (test_labels - mean_vals_y) / std_devs_y
# padding test features with 1

ones = np.ones((test_features.shape[0], 1))
test_features_padded = np.hstack((test_features, ones))
```

```

linear_pred_train = X_padded @ B_linear
ridge_pred_train = X_padded @ B_ridge

linear_pred_test = test_features_padded @ B_linear
ridge_pred_test = test_features_padded @ B_ridge

linear_train_error_2 = ((y - linear_pred_train).T @ (y - linear_pred_train)) / (2 * (X_padded.shape[0]))
ridge_train_error_2 = ((y - ridge_pred_train).T @ (y - ridge_pred_train)) / (2 * (X_padded.shape[0]))

linear_test_error_2 = ((test_labels - linear_pred_test).T @ (test_labels - linear_pred_test)) / (2 * (test_features_padded.shape[0]))
ridge_test_error_2 = ((test_labels - ridge_pred_test).T @ (test_labels - ridge_pred_test)) / (2 * (test_features_padded.shape[0]))


print("Linear Regression Training Error:", linear_train_error_2)
print("Ridge Regression Training Error:", ridge_train_error_2)
print("Linear Regression Testing Error:", linear_test_error_2)
print("Ridge Regression Testing Error:", ridge_test_error_2)

```

```

Linear Regression Training Error: 0.049144184358431243
Ridge Regression Training Error: 0.04914418603031624
Linear Regression Testing Error: 0.04739264015617485
Ridge Regression Testing Error: 0.04739197713322119

```

```

import matplotlib.pyplot as plt
import numpy as np

error_types = ['Linear Regression Training', 'Ridge Regression Training', 'Linear Regression Testing', 'Ridge Regression Test']
num_error_types = len(error_types)

cases = ['Case 1', 'Case 2']
error_values = np.array([
    [linear_train_error_1, ridge_train_error_1, linear_test_error_1, ridge_test_error_1],
    [linear_train_error_2, ridge_train_error_2, linear_test_error_2, ridge_test_error_2]
])

# Define colors for the bars for each case
colors = ['r', 'y'] # Blue for Case 1, Green for Case 2

bar_width = 0.35

x = np.arange(num_error_types)

fig, ax = plt.subplots(figsize=(10, 6))

for i, case in enumerate(cases):
    ax.bar(x + i * bar_width, error_values[i], bar_width, label=case, color=colors[i])

ax.set_xlabel('Error Types')
ax.set_ylabel('Error Value')
ax.set_title('Comparison of Errors for case 1 (didnt normalise/stnd) and Case 2 ( noramlize / stnd)')
ax.set_xticks(x + bar_width / 2)
ax.set_xticklabels(error_types)
ax.legend()

plt.tight_layout()
plt.show()

```



▼ 2) Logistic Regression

0.8 +

$$\frac{\partial}{\partial \theta_j} J(\theta) = \frac{1}{m} \sum_{i=1}^m \left(h_\theta(x^{(i)}) - y^{(i)} \right) x_j^{(i)}$$

```
data = pd.read_csv("/content/drive/My Drive/ISL/banknote_authentication.csv")
data.head(5)
```

	variance	skewness	curtosis	entropy	class
0	3.62160	8.6661	-2.8073	-0.44699	0
1	4.54590	8.1674	-2.4586	-1.46210	0
2	3.86600	-2.6383	1.9242	0.10645	0
3	3.45660	9.5228	-4.0112	-3.59440	0
4	0.32924	-4.4552	4.5718	-0.98880	0

```
data = data.values
```

split ratio = 0.8

```
total_samples = len(data)
train_samples = int(total_samples * split_ratio)
test_samples = total_samples - train_samples
```

```
train_data = data[:train_samples]
test_data = data[train_samples:]
```

```
X = train_data[ : , : -1]  
y = train_data[ : , -1]
```

```
test_X = test_data[ :, : -1]  
test_y = test_data[ :, -1]
```

```
# padding the train data  
ones = np.ones((X.shape[0] , 1))  
X_padded = np.hstack((X , ones))  
X_padded
```

```
array([[ 3.6216 ,  8.6661 , -2.8073 , -0.44699,  1.      ],
       [ 4.5459 ,  8.1674 , -2.4586 , -1.4621 ,  1.      ],
       [ 3.866 , -2.6383 ,  1.9242 ,  0.10645,  1.      ],
       ...,
       [ 2.0177 ,  1.7982 , -2.9581 ,  0.2099 ,  1.      ],
       [ 1.164 ,  3.913 , -4.5544 , -3.8672 ,  1.      ],
       [-4.3667 ,  6.0692 ,  0.57208, -5.4668 ,  1.      ]])
```

```
def sigm( x , theta):  
    z = np.dot( x , theta)  
    ep = np.exp(-z)  
    return ( 1 ) / ( 1 + ep )
```

```

# padding test-X
ones = np.ones((test_X.shape[0] , 1))
test_X_padded = np.hstack((test_X ,ones))

theta = [0 , 0 ,0 ,0 ,0 ]

# gradient descent for the optimal parameters
alpha = 0.5
costs = []

iters = 1000

for i in range(iters):
    ptheta = theta
    for j in range(len(theta)):
        temp = 0;
        for xx, yy in zip(X_padded , y ) :
            fx = sigm(xx, ptheta)
            temp = temp + (( fx - yy ) * (xx[j]))
        temp /= (X_padded.shape[0])
        temp *= alpha
        theta[j] -= temp

    st = X_padded @ theta
    prob = ( (1) / ( 1 + np.exp(-st)) )

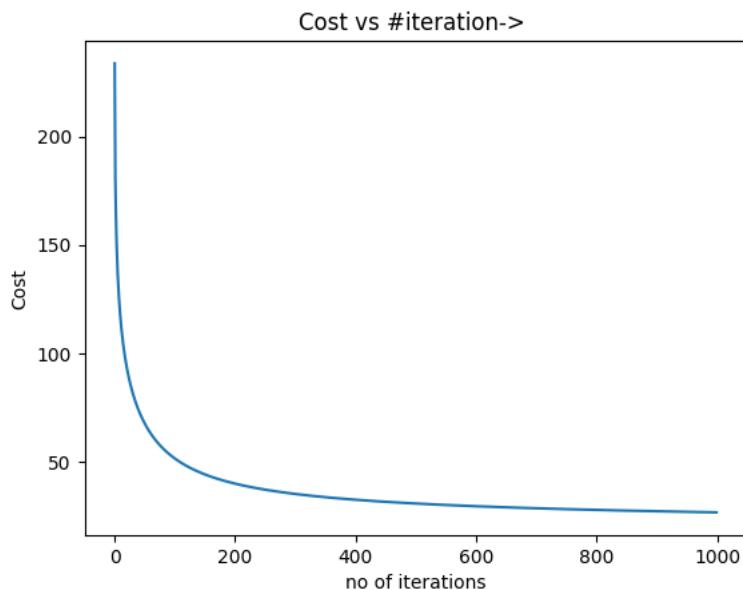
    cost = 0
    for predd , yy in zip( prob , y ):
        if yy == 1:
            cost += -(np.log(predd))
        else:
            cost += -(np.log(1- predd))
    costs.append(cost)

```

```

plt.plot(costs)
plt.title("Cost vs #iteration-> ")
plt.xlabel("no of iterations")
plt.ylabel("Cost")
plt.show()

```



```
theta
```

```
[ -3.426170739063408,
 -1.8817312945138627,
```

```
-2.3006951893166594,
-0.09374573346865879,
3.214515797551809]
```

```
st = test_X_padded @ theta

prob = ( (1) /( 1 + np.exp(-st)) )
pred = ( prob >= 0.5).astype(int)

test_cost = 0 ;
for predd , yy in zip( prob , test_y ):
    if yy == 1:
        test_cost += -(np.log(predd))
    else:
        test_cost += -(np.log(1- predd))

print(test_cost)
```

7.765000364905097

```
# defining accuracy as correctly identified / total test
```

```
num = ( pred == test_y)
# num = np.array(num)

accuracy = np.count_nonzero(num) / (test_y.shape[0])
print(f"{accuracy} is the test accuracy")
```

```
st = X_padded @ theta

prob = ( (1) /( 1 + np.exp(-st)) )
pred_train = ( prob >= 0.5).astype(int)

num = ( pred_train == y)

accuracy = np.count_nonzero(num) / (y.shape[0])
print(f"{accuracy} is the train accuracy")
```

```
train_cost = 0 ;
for predd , yy in zip( prob , y ):
    if yy == 1:
        train_cost += -(np.log(predd))
    else:
        train_cost += -(np.log(1- predd))
```

```
print()
print(f"{train_cost} is the training cost" )
print(f"{test_cost} is the testing cost" )
```

0.9927272727272727 is the test accuracy
0.9863263445761167 is the train accuracy

26.7434716825555 is the training cost
7.765000364905097 is the testing cost

As we can see the accuracy that we are getting is > 98 %

also the train error and train error are comparable hence it is a good and generalized model

▼ 3) Understanding Bias Variance Trade off

- ▼ 1. Load the data and plot the stock value Vs time.

```
data = pd.read_excel("/content/drive/My Drive/ISL/Apple_stock_data.xlsx")
data.head(5)
```

	Date	Close/Last	Volume	Open	High	Low	grid icon
0	09/29/2023	\$171.21	51861080	\$172.02	\$173.07	\$170.341	grid icon
1	09/28/2023	\$170.69	56294420	\$169.34	\$172.03	\$167.62	grid icon
2	09/27/2023	\$170.43	66921810	\$172.62	\$173.04	\$169.05	grid icon
3	09/26/2023	\$171.96	64588950	\$174.82	\$175.20	\$171.66	grid icon
4	09/25/2023	\$176.08	46172740	\$174.20	\$176.97	\$174.15	grid icon

```
df = data
df['Close/Last'].dtype
```

```
dtype('O')
```

```
df['Date'] = pd.to_datetime(df['Date'])
```

```
df['Close/Last'] = df['Close/Last'].str.replace('$', '').astype(float)
```

```
<ipython-input-296-b53857874181>:3: FutureWarning: The default value of regex will change from True to False in a future
df['Close/Last'] = df['Close/Last'].str.replace('$', '').astype(float)
```

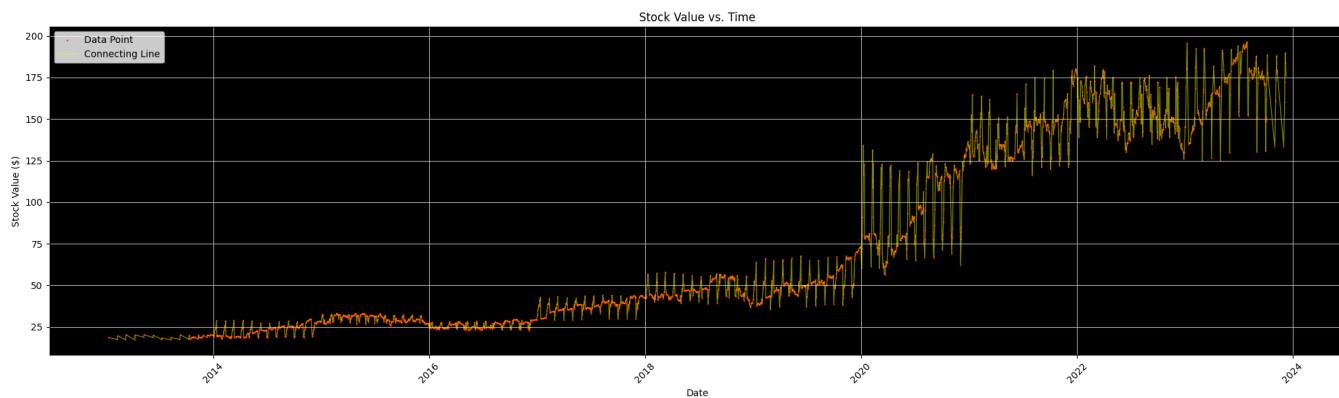
```
df.sort_values(by='Date', inplace=True)
```

```
plt.figure(figsize=(20, 6))
plt.gca().set_facecolor('black')
plt.scatter(df['Date'], df['Close/Last'], marker='o', label='Data Point', color='red' , s = 0.3)

plt.plot(df['Date'], df['Close/Last'], linestyle='-', linewidth=1, color='yellow', alpha=0.5, label='Connecting Line')

plt.title('Stock Value vs. Time')
plt.xlabel('Date')
plt.ylabel('Stock Value ($)')
plt.grid(True)
plt.xticks(rotation=45)
plt.tight_layout()
plt.legend()

plt.show()
```



2. Split the data into train and test. (data upto 2022, consider as training data, remaining data can be used as testing data)

```
train_data = df[df['Date'] <= '2022-12-31'] # Data up to the end of 2022
test_data = df[df['Date'] > '2022-12-31'] # Data from 2023 onwards
```

```
train_data.shape
```

```
(2329, 6)
```

```
train_data.reset_index(drop=True, inplace=True)
test_data.reset_index(drop=True, inplace=True)
```

```
train_data.shape
```

```
(2329, 6)
```

```
df.head(5)
```

	Date	Close/Last	Volume	Open	High	Low	grid
2493	2013-01-11	18.5725	274097882	\$18.715	\$18.7429	\$18.4229	grid
2515	2013-02-10	17.4843	287196445	\$17.3439	\$17.5643	\$17.2768	grid
2473	2013-02-12	19.6868	471672691	\$19.9286	\$20.1546	\$19.6722	grid
2514	2013-03-10	17.2646	315209314	\$17.5182	\$17.5839	\$17.1692	grid
2472	2013-03-12	20.2258	449352780	\$19.9393	\$20.2279	\$19.9172	grid

3. Apply polynomial fit to predict the current stock value from the past values.

$$Y_t = \beta_{1,t-1} Y_{t-1} + \beta_{2,t-1} Y_{t-1}^2 + \dots + \beta_{p-1,t-k} Y_{t-k}^{p-1} + \beta_{p,t-k} Y_{t-k}^p$$

4. Add more features using the higher powers of the past stock values in the dataset. We have two parameters in hand i) k- past window length ii)p-order of the polynomial. Play with $k = 1, 2, 3$ and $p = 1, 2, 3$. k and p should be atleast 3.
5. Computer the prediction accuracy for the test data and plot the predicted stock value and original stock value Vs time on a single plot. (For all values of k and p). i) Observe the change in parameter values β w.r.t p and k ii) comment on bias and variance.

```
def plotfig( y , pred , k , p ):
    plt.figure(figsize=(10, 6))
    plt.plot(y, label='Original Stock Values', color='blue')

    plt.plot(pred, label=f'Predicted Values (k={k}, p={p})', color='red')

    plt.xlabel('Time')
    plt.ylabel('Stock Value')
    plt.title(f'Predicted vs. Original Stock Values (k={k}, p={p})')
    plt.legend()
    plt.show()
```

```
import numpy as np
from sklearn.metrics import mean_squared_error
import matplotlib.pyplot as plt
```

```
# Define ks and ps
```

```

ks = [3, 4, 5, 7, 9, 10]
ps = [3, 4, 5, 6, 8, 10]

# taking three cases
# k = 3 , p = 3 i.e high bias
# k = 5 , p= 5 i.e. generalized
# k = 10 , p = 10 , i.e. high variance

k3p3_test = np.array([])
k3p3_train = np.array([])
k5p5_test = np.array([])
k5p5_train = np.array([])
k10p10_test = np.array([])
k10p10_train = np.array([])

te = []
tse = []

pred_error_p = []
pred_error_k = []

thetas = [ [] ]

test_errors = []
train_errors = []

for k in ks:
    for p in ps:

        X_train = np.zeros((len(train_data) - k, p))
        y_train = train_data['Close/Last'][k:]

        X_test = np.zeros((len(test_data) - k, p))
        y_test = test_data['Close/Last'][k:]

        for kk in range(k):
            for pp in range(p):
                X_train[:, pp] = train_data['Close/Last'][k - kk - 1:-kk - 1] ** (pp + 1)
                X_test[:, pp] = test_data['Close/Last'][k - kk - 1:-kk - 1] ** (pp + 1)

        # Fit the model
        theta = np.linalg.inv(X_train.T @ X_train) @ (X_train.T @ y_train)

        pred_train = X_train @ theta
        pred_test = X_test @ theta

        # errors
        squared_diff_train = (y_train - pred_train) ** 2
        train_error = np.mean(squared_diff_train)

        squared_diff_test = (y_test - pred_test) ** 2
        test_error = np.mean(squared_diff_test)

        train_errors.append(train_error)
        test_errors.append(test_error)

        print(f"k = {k}, p = {p}")
        print(f"Training error: {train_error}")
        print(f"Testing error: {test_error}")

        plotfig(y_test, pred_test, k, p)
        plotfig(y_train, pred_train, k, p)

        if k == 3 and p == 3:
            k3p3_test = pred_test
            k3p3_train = pred_train
            te.append(train_error)
            tse.append(test_error)

        if k == 5 and p == 5:
            k5p5_test = pred_test
            k5p5_train = pred_train
            te.append(train_error)
            tse.append(test_error)

        if k == 10 and p == 10:
            k10p10_test = pred_test
            k10p10_train = pred_train
            te.append(train_error)
            tse.append(test_error)

```

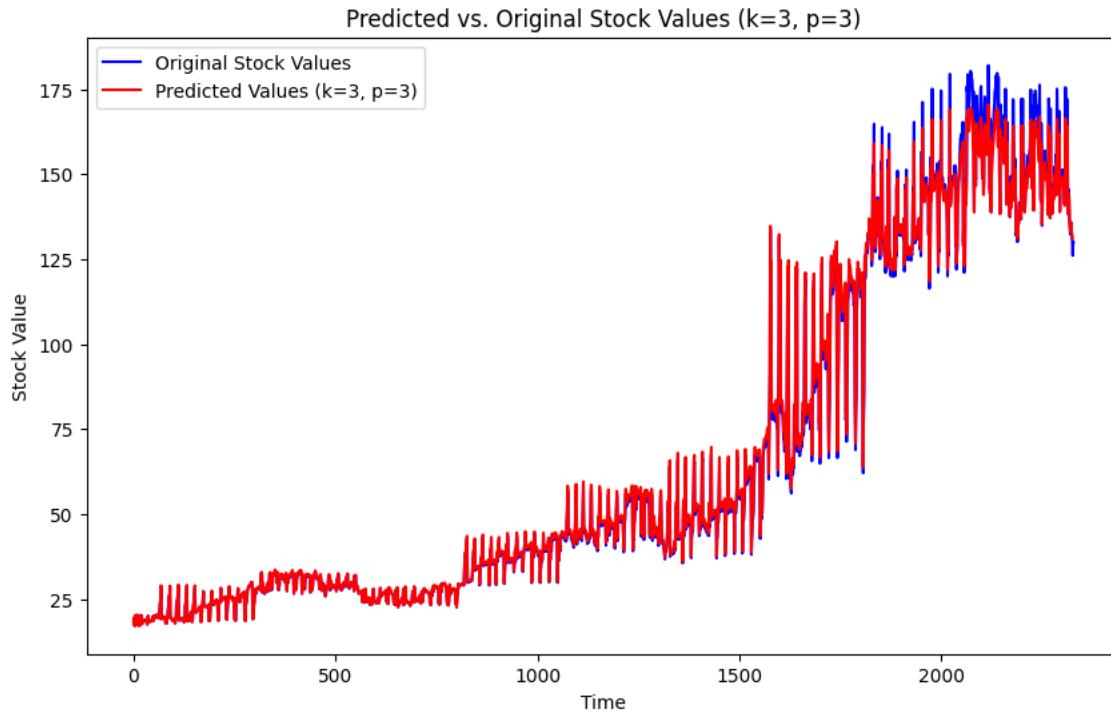
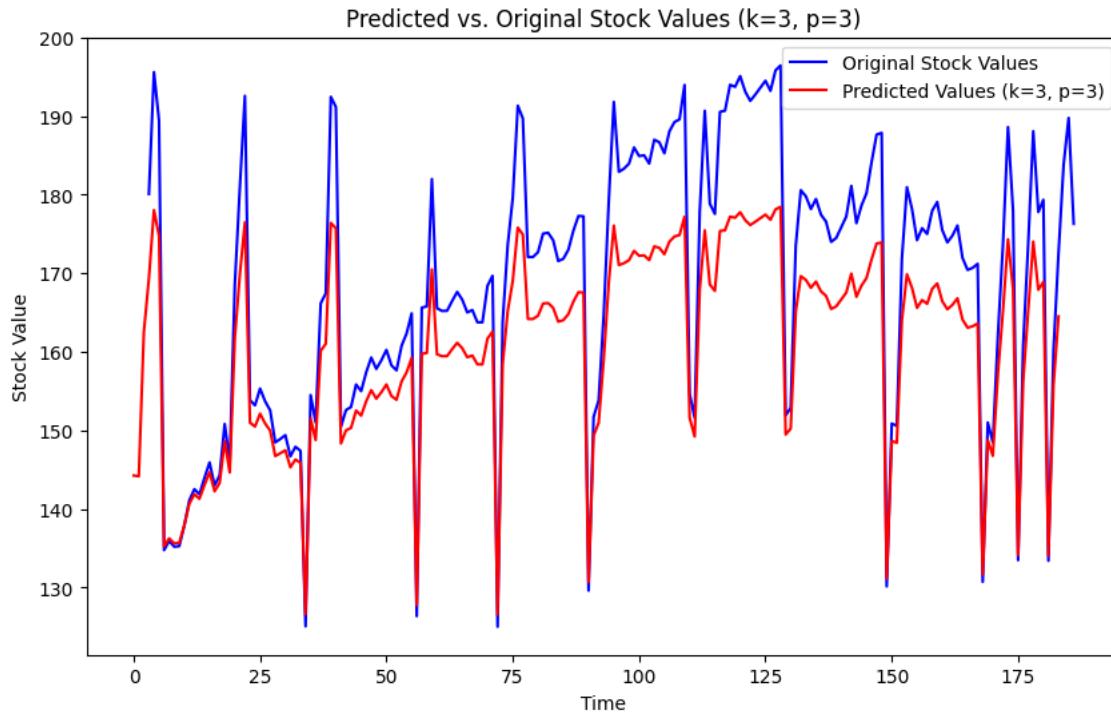
```
k10p10_test = pred_test
k10p10_train = pred_train
te.append(train_error)
tse.append(test_error)

thetas.append(theta)

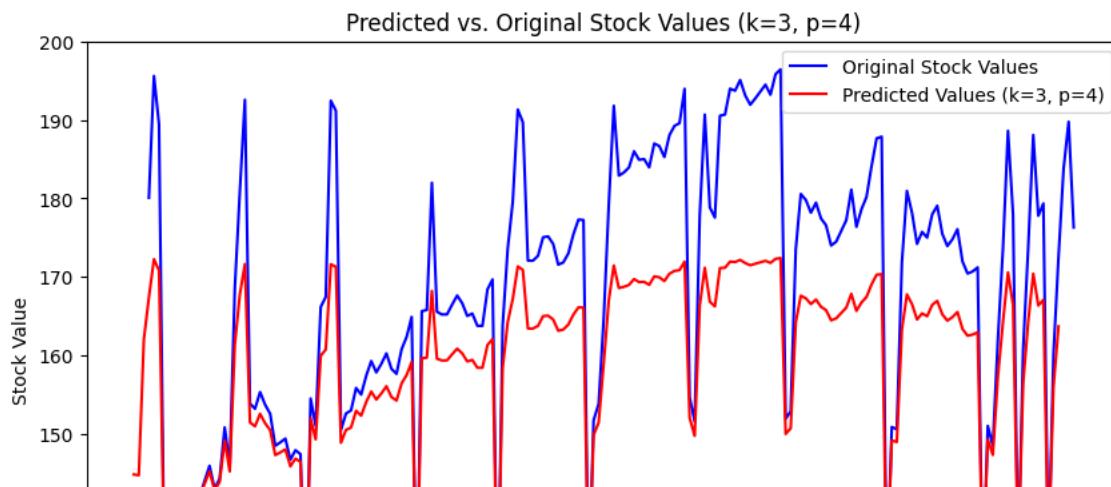
if k == 5:
    pred_error_p.append(test_error)

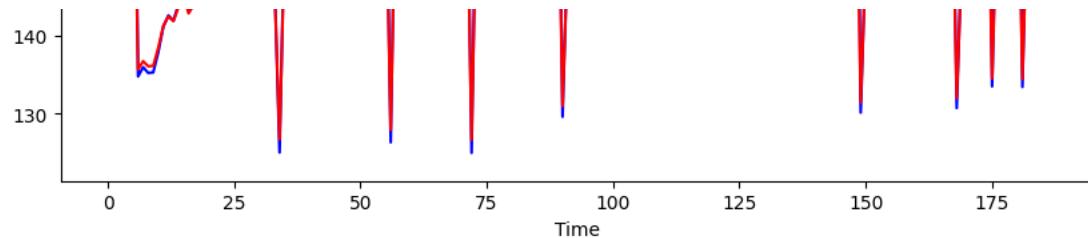
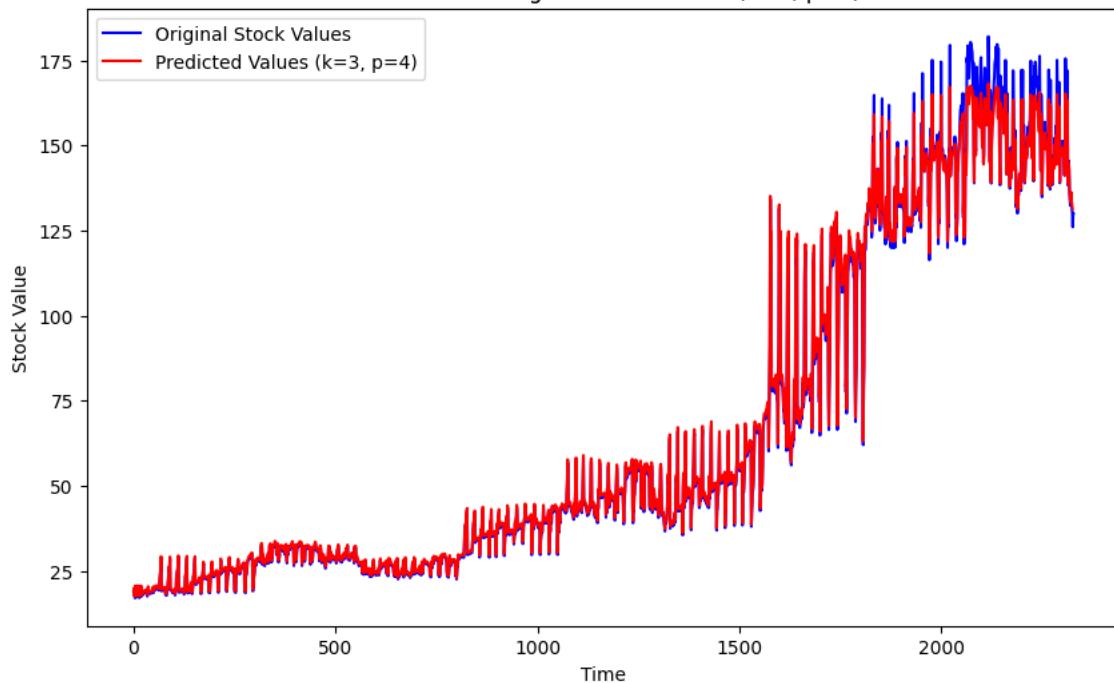
if p == 5:
    pred_error_k.append(test_error)
```

$k = 3, p = 3$
 Training error: 104.35061822387826
 Testing error: 418.2981912312131



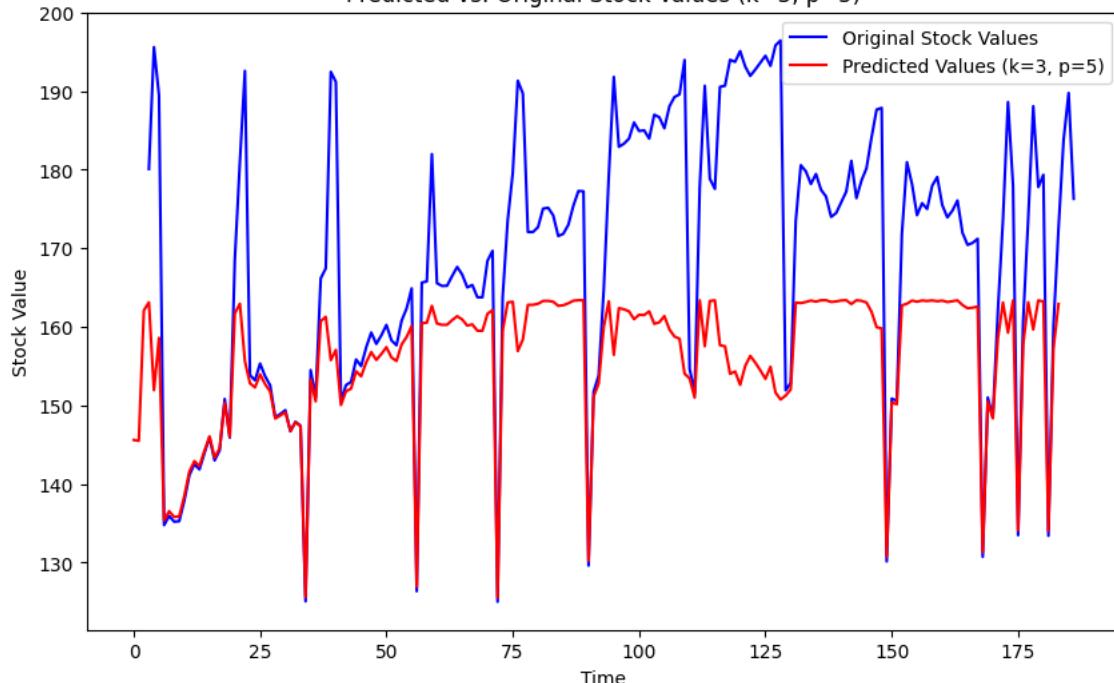
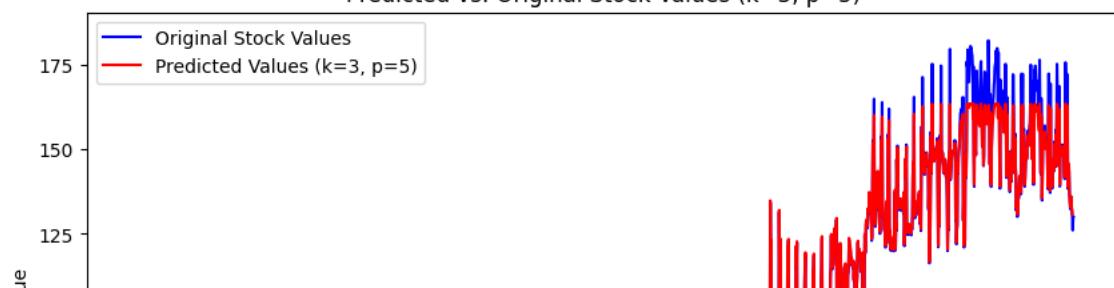
$k = 3, p = 4$
 Training error: 104.156439182689
 Testing error: 419.93692937472593

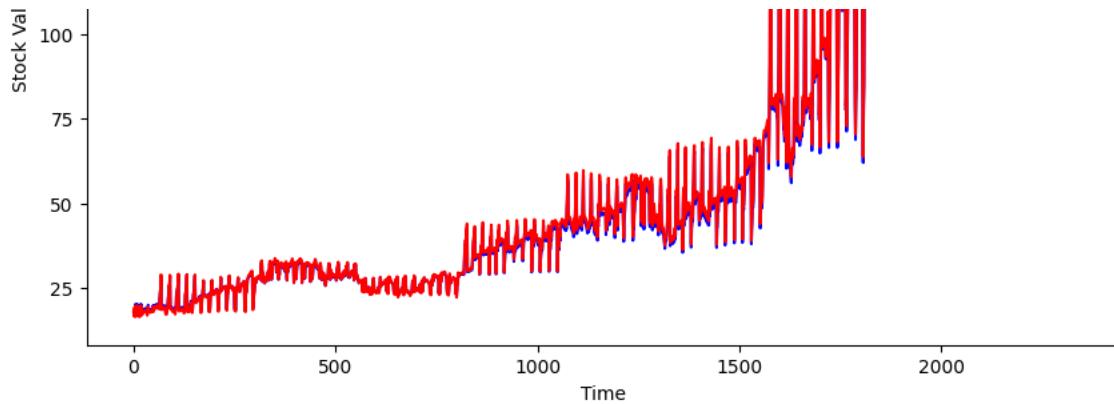


Predicted vs. Original Stock Values ($k=3, p=4$) $k = 3, p = 5$

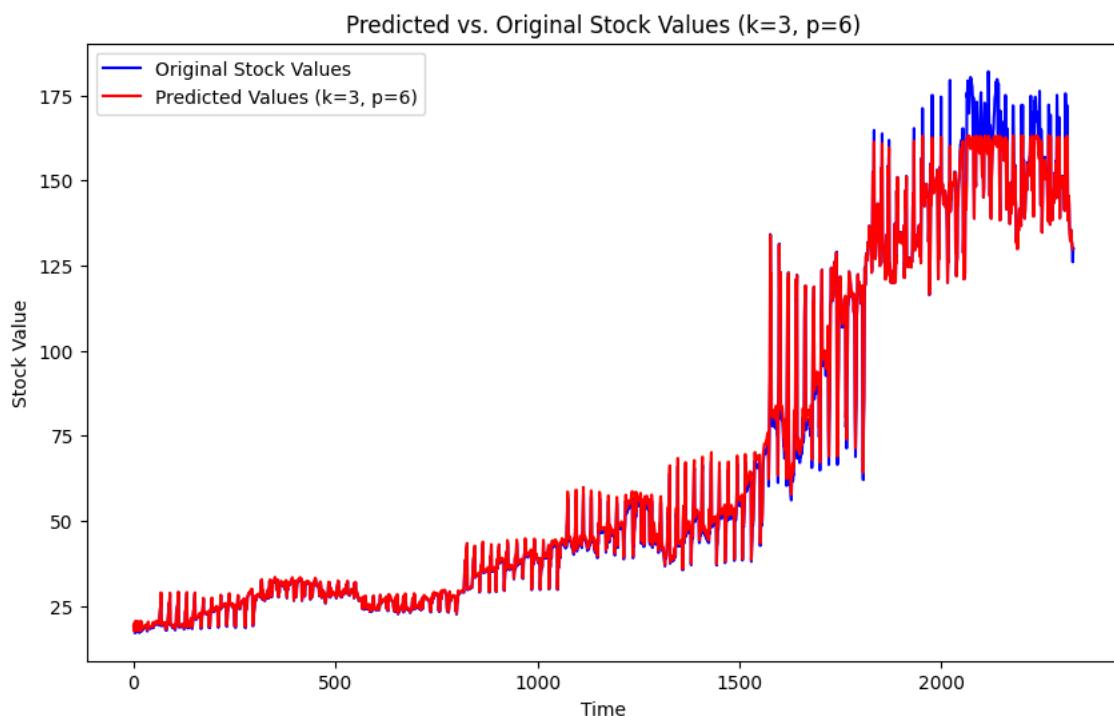
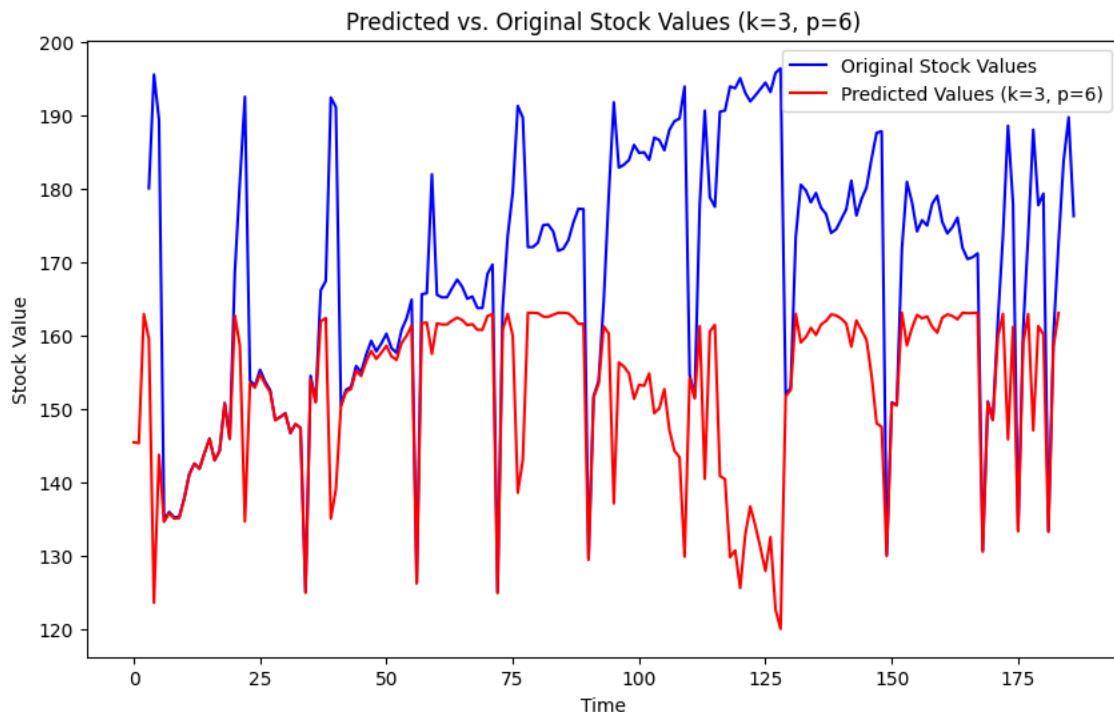
Training error: 103.50715707006074

Testing error: 485.1309157313537

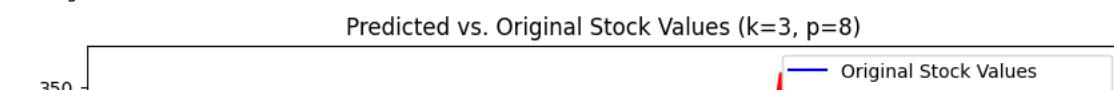
Predicted vs. Original Stock Values ($k=3, p=5$)Predicted vs. Original Stock Values ($k=3, p=5$)

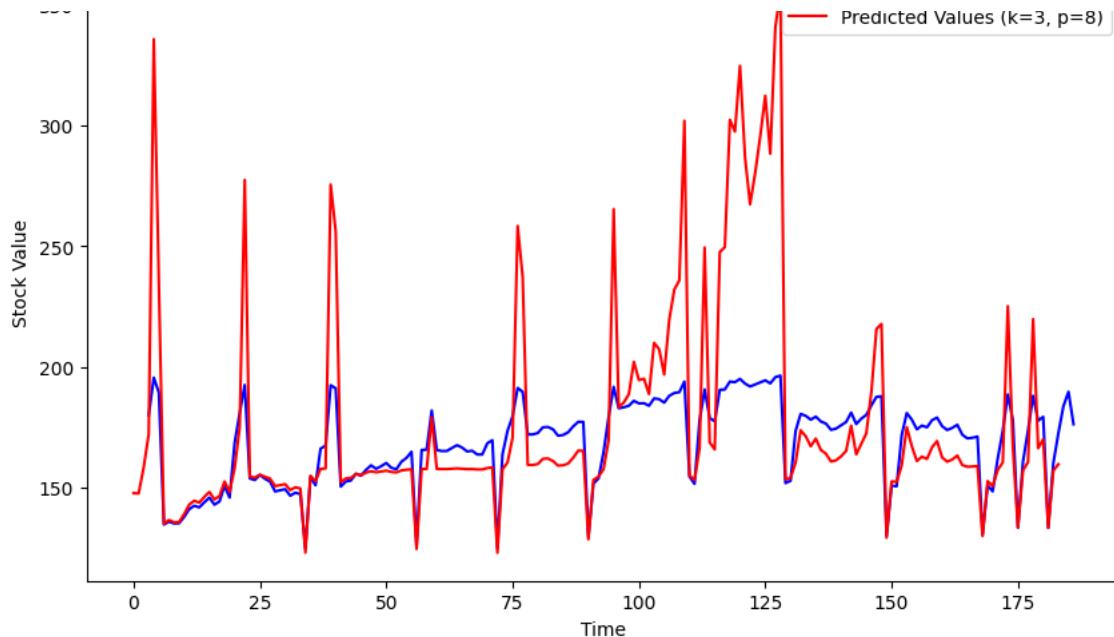


$k = 3, p = 6$
 Training error: 103.1004733736073
 Testing error: 699.3396580141628

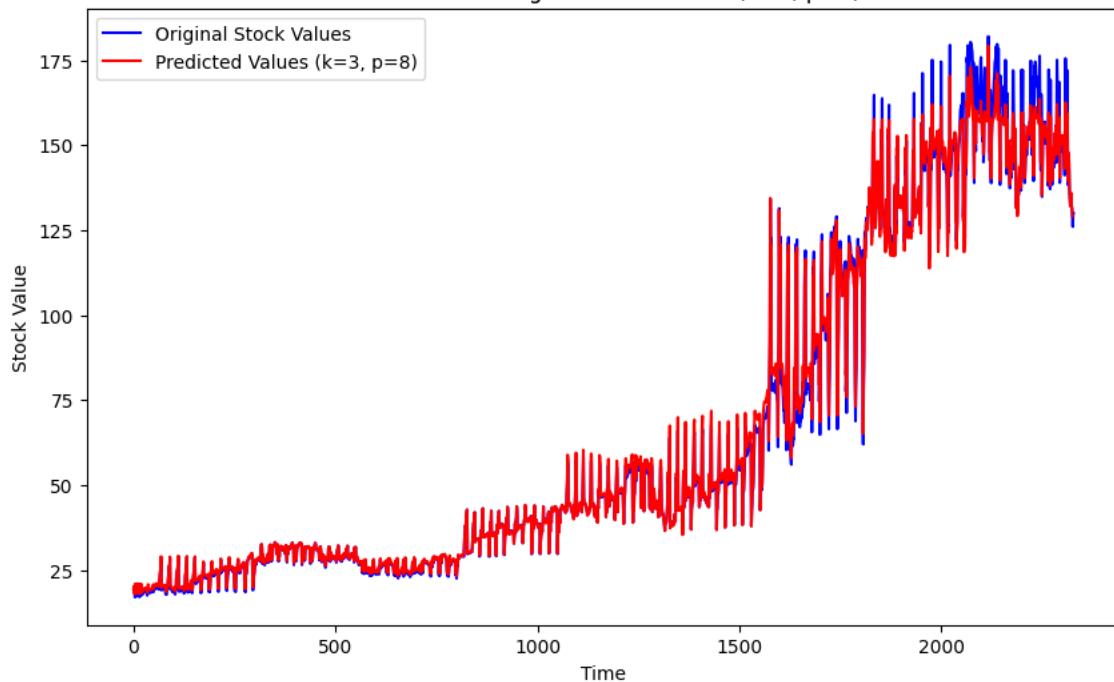


$k = 3, p = 8$
 Training error: 100.90855893819274
 Testing error: 2200.7199025502746





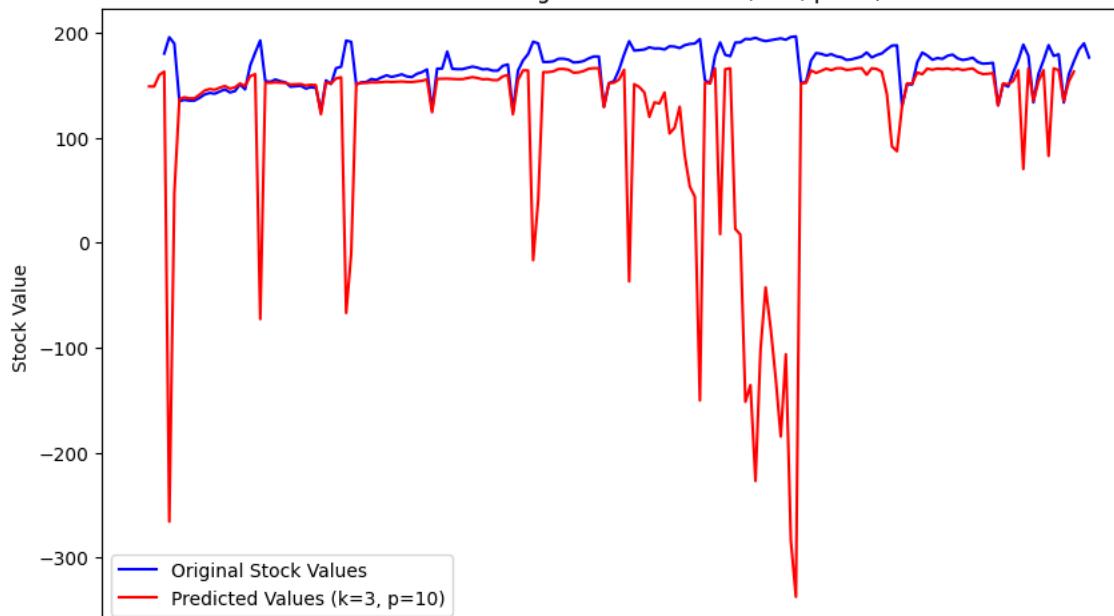
Predicted vs. Original Stock Values (k=3, p=8)

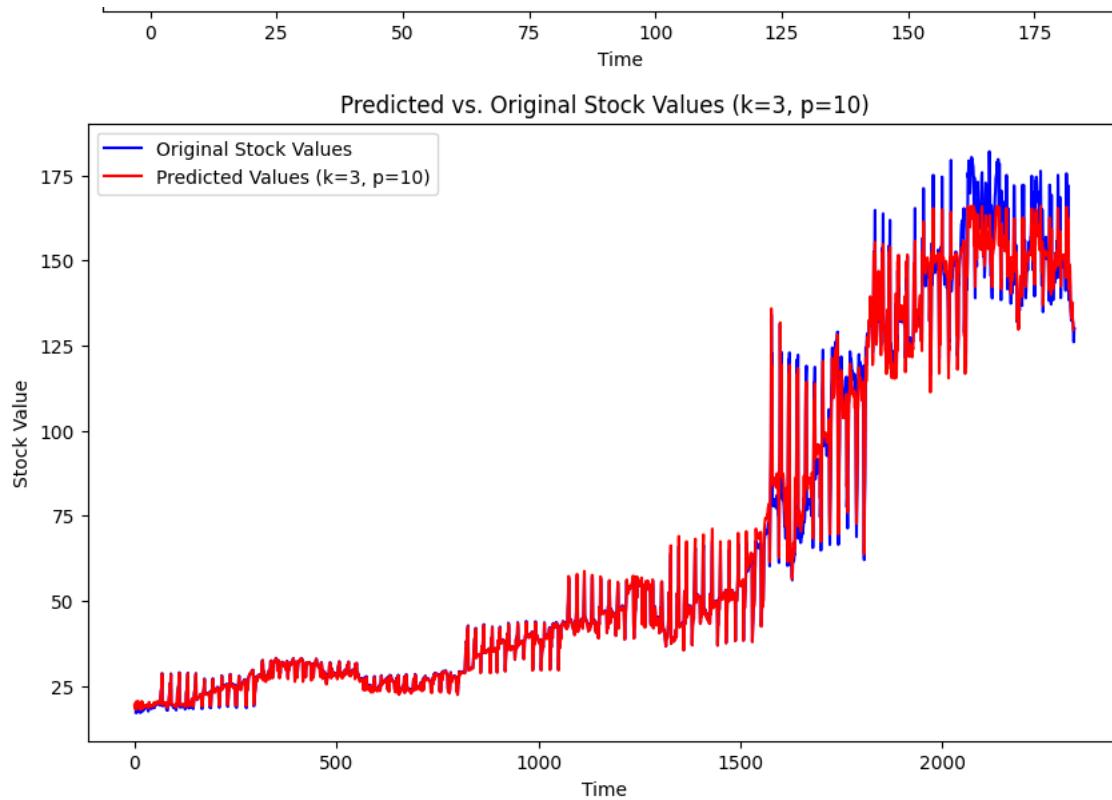
 $k = 3, p = 10$

Training error: 99.56244740681957

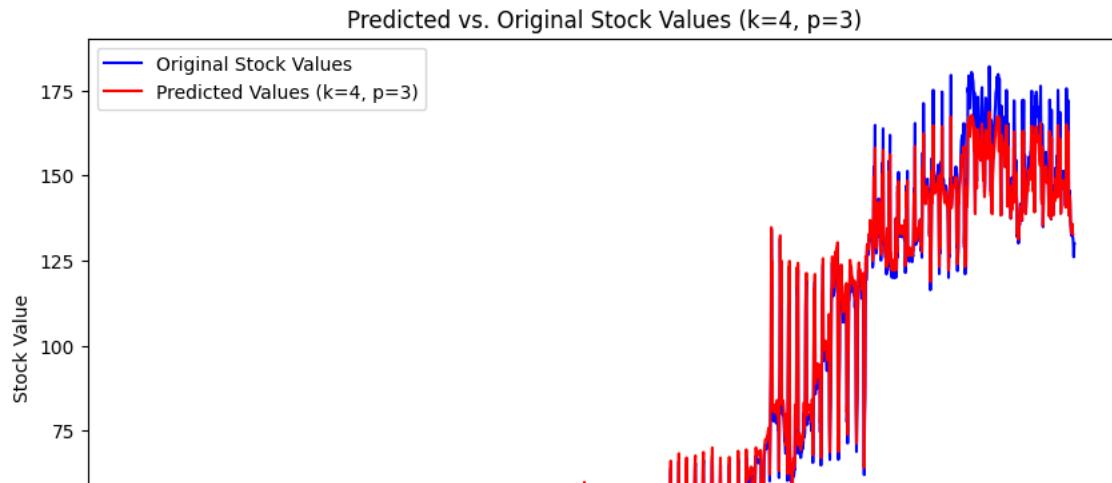
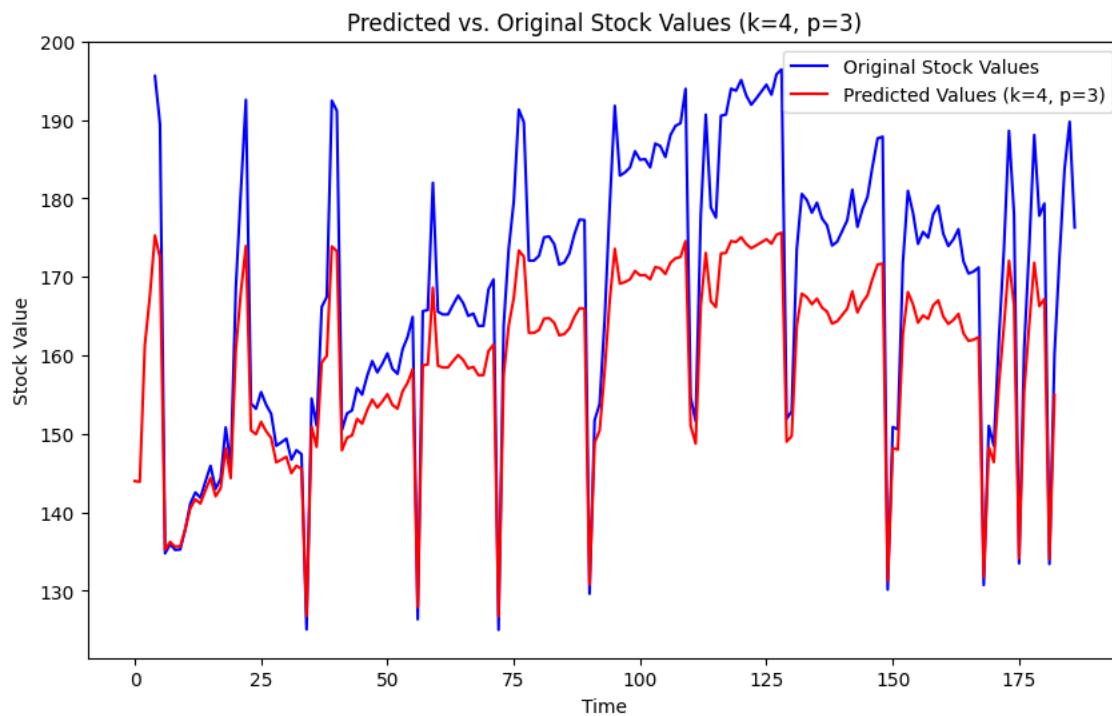
Testing error: 11659.97826341857

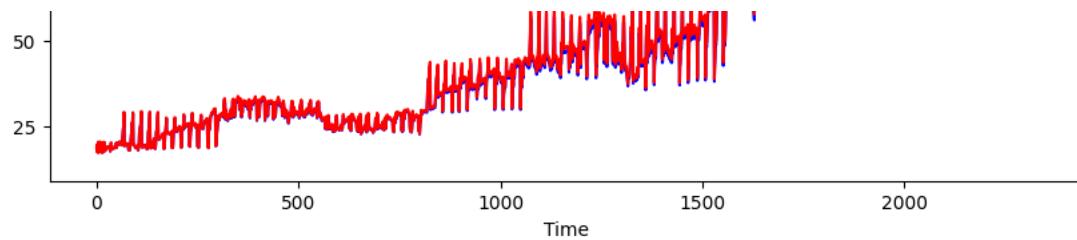
Predicted vs. Original Stock Values (k=3, p=10)



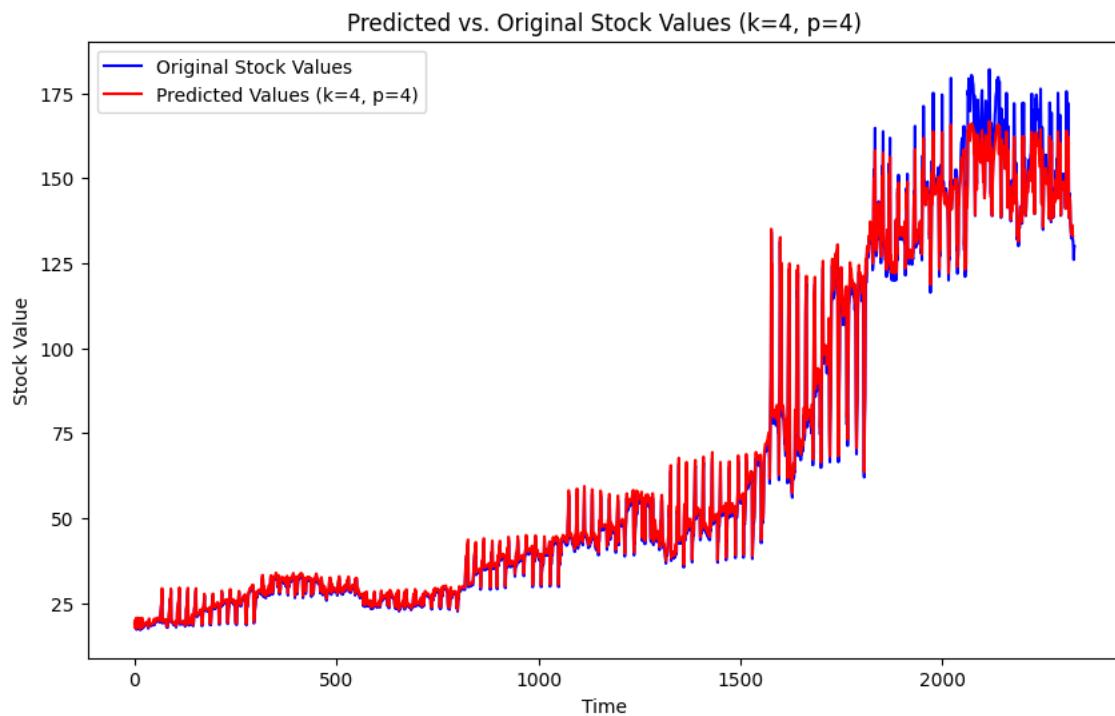
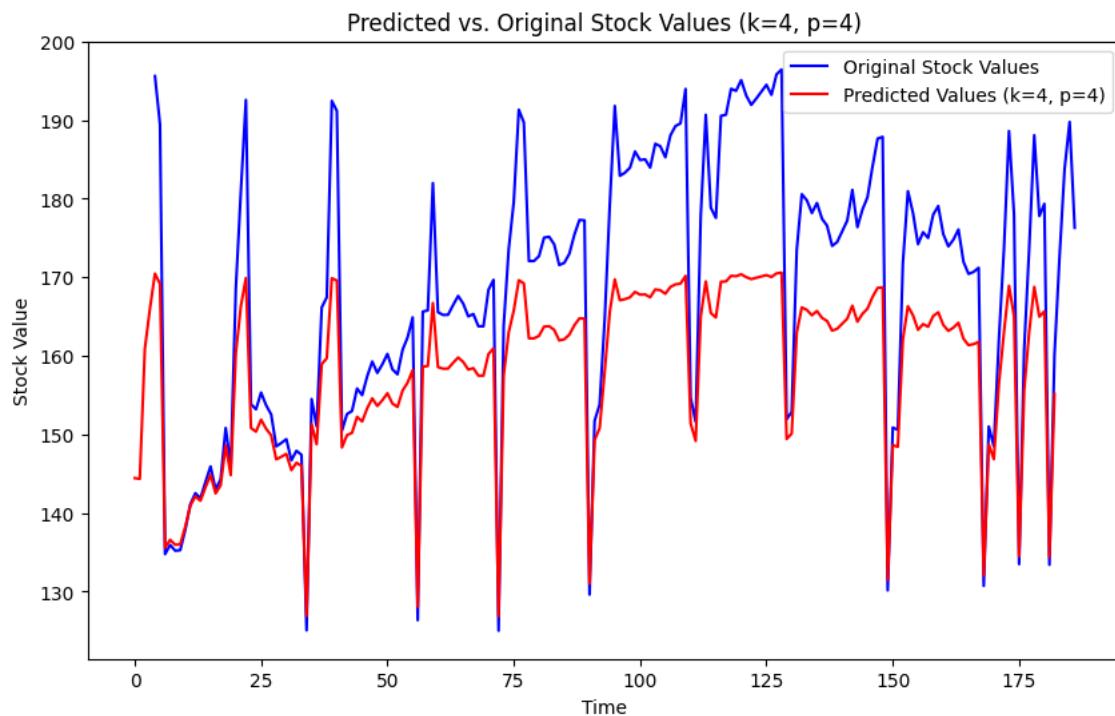


$k = 4, p = 3$
Training error: 125.49687651021092
Testing error: 437.69698368342637

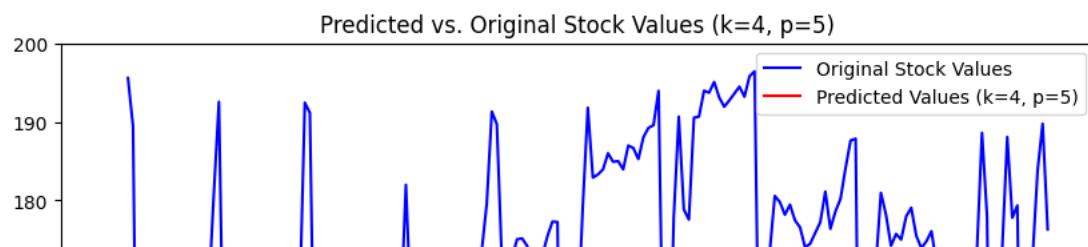


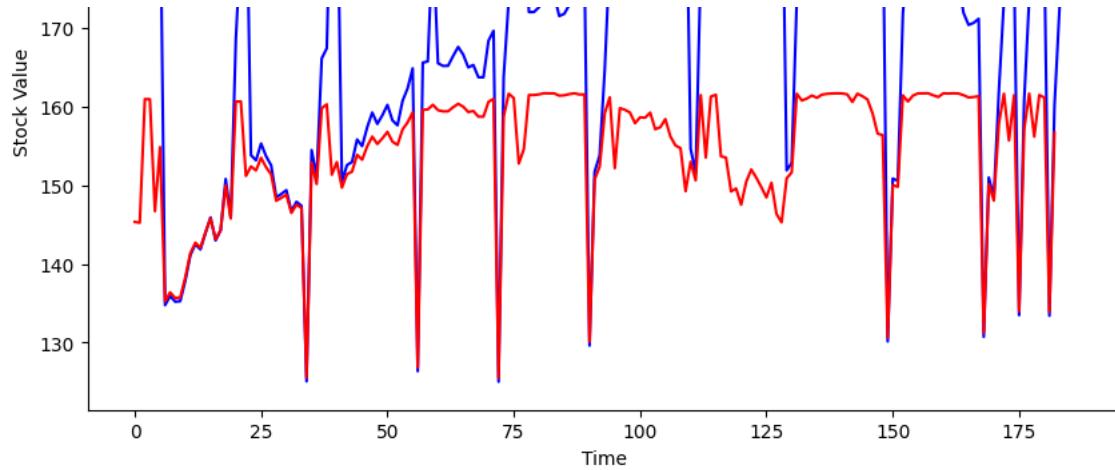
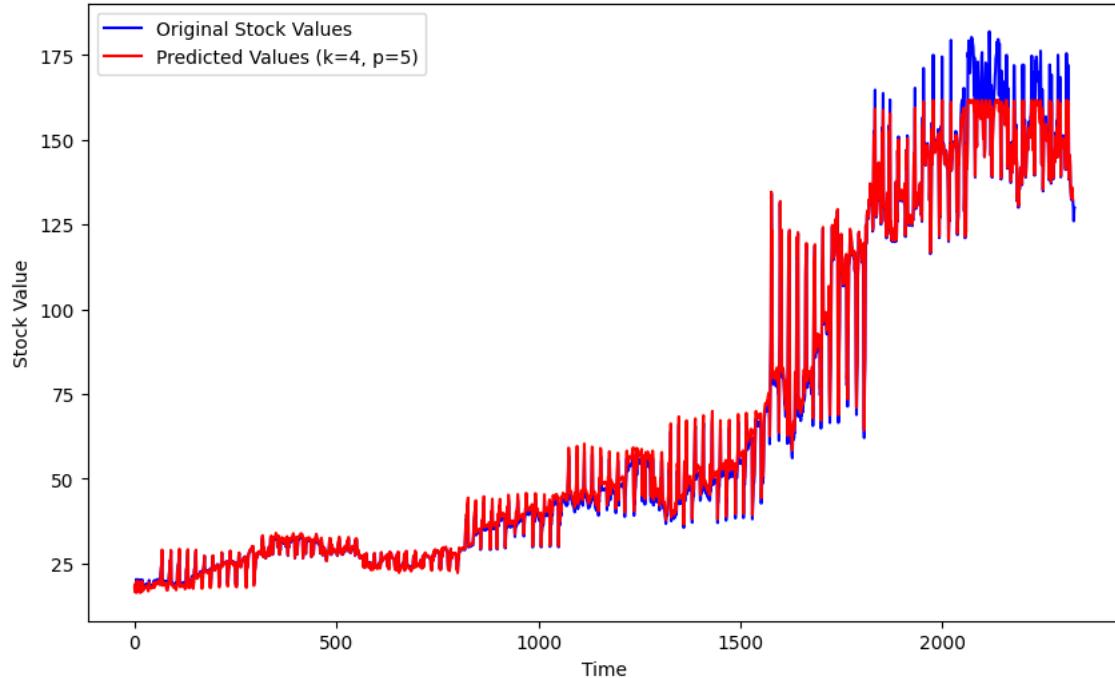


$k = 4, p = 4$
 Training error: 125.36182581561926
 Testing error: 444.9018243045056

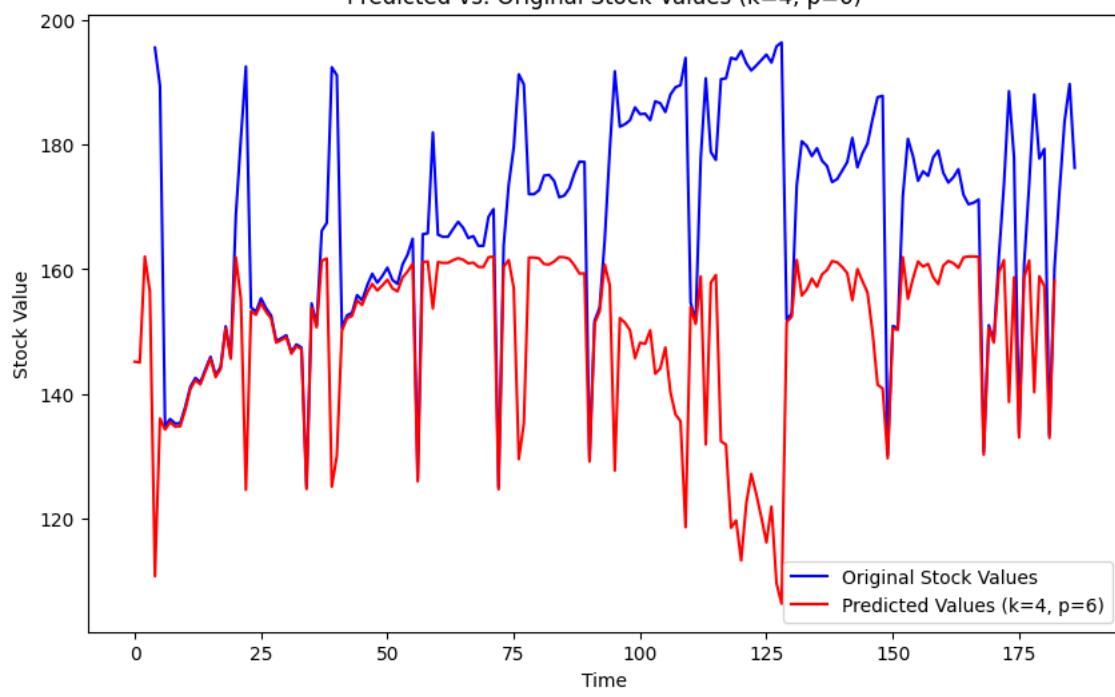


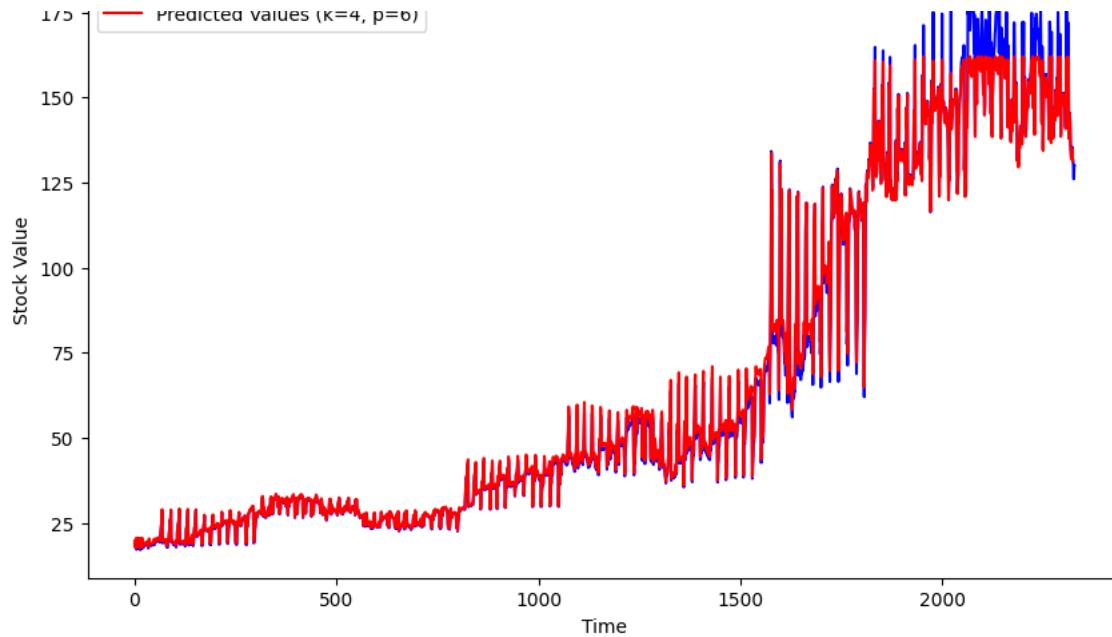
$k = 4, p = 5$
 Training error: 124.47510591703555
 Testing error: 555.075876148586



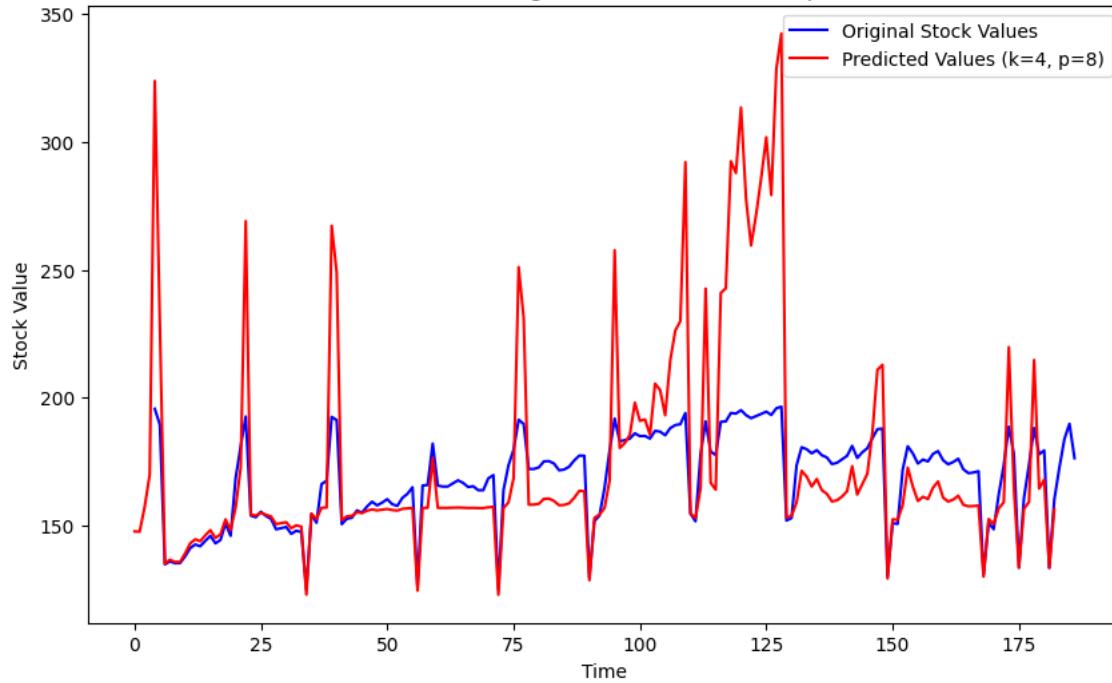
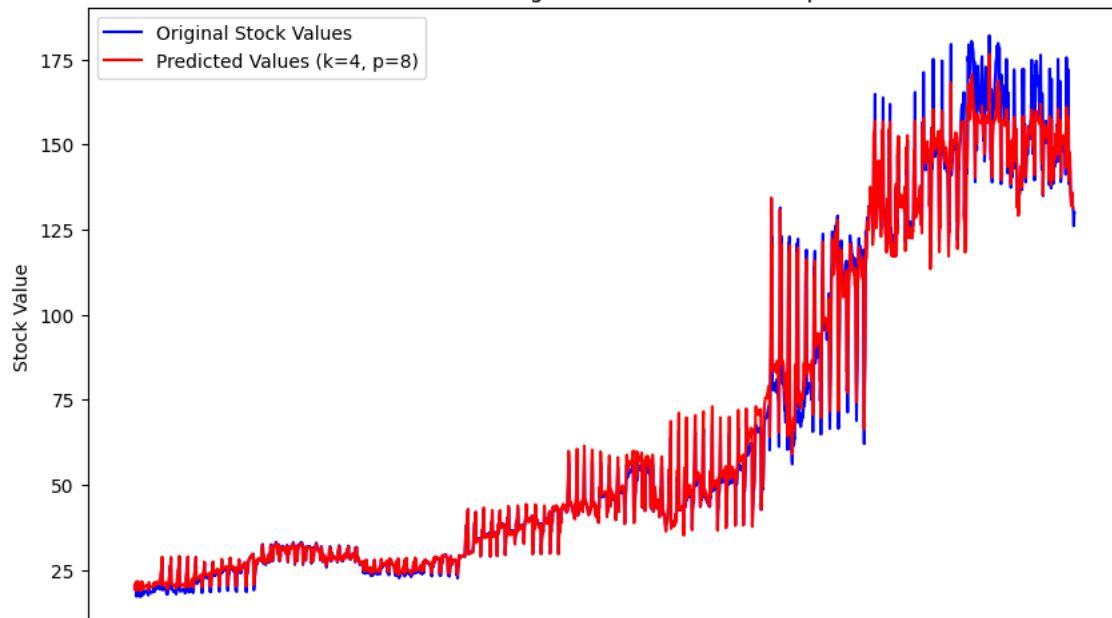
Predicted vs. Original Stock Values ($k=4, p=5$)

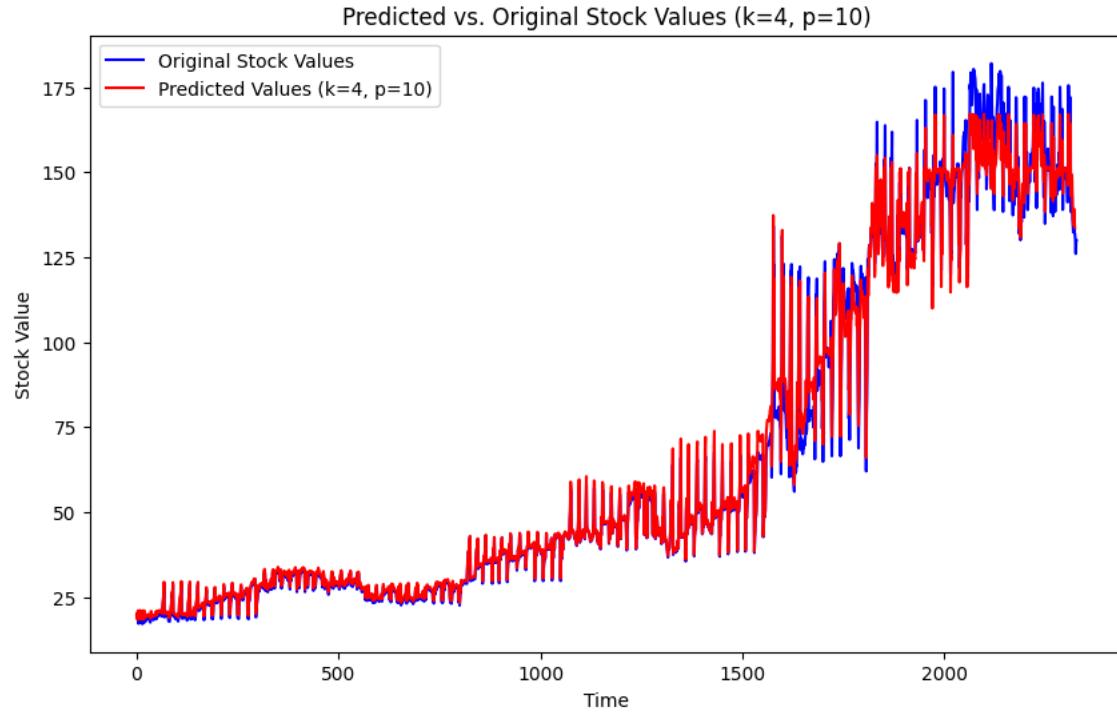
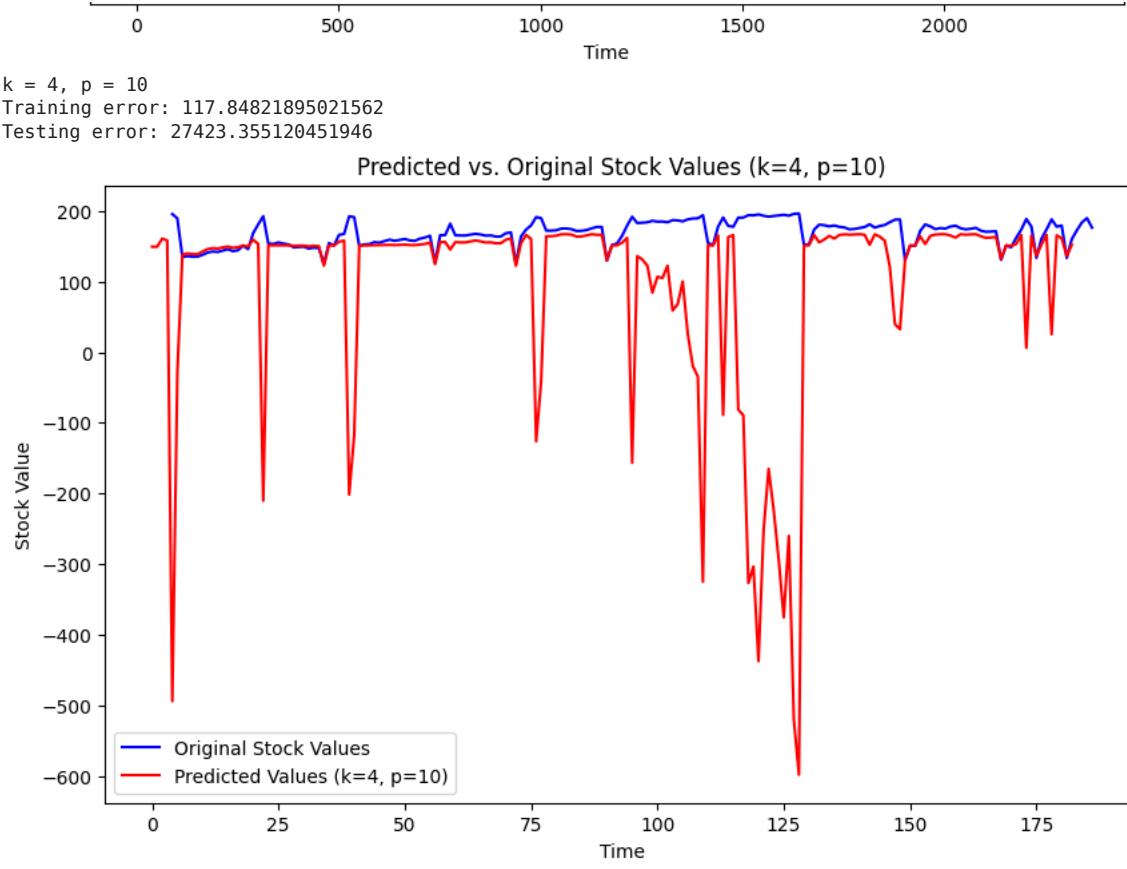
$k = 4, p = 6$
Training error: 123.8199834853926
Testing error: 912.1810208935289

Predicted vs. Original Stock Values ($k=4, p=6$)Predicted vs. Original Stock Values ($k=4, p=6$)

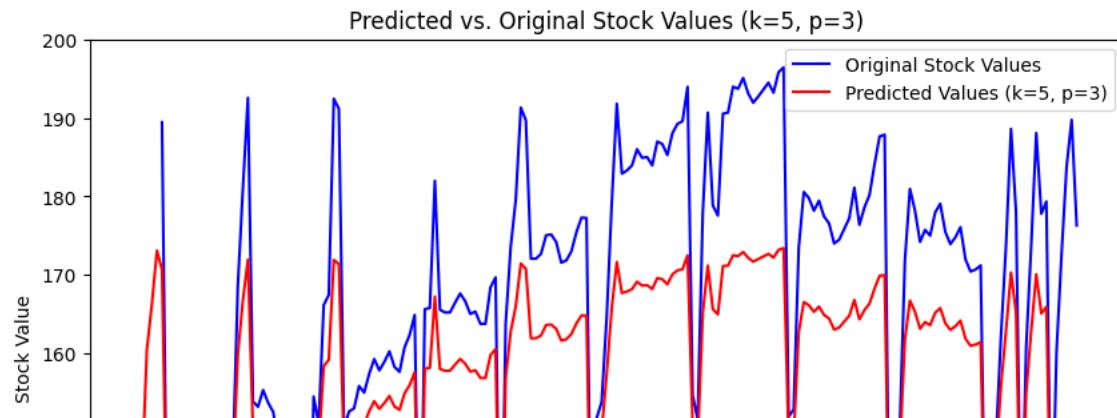


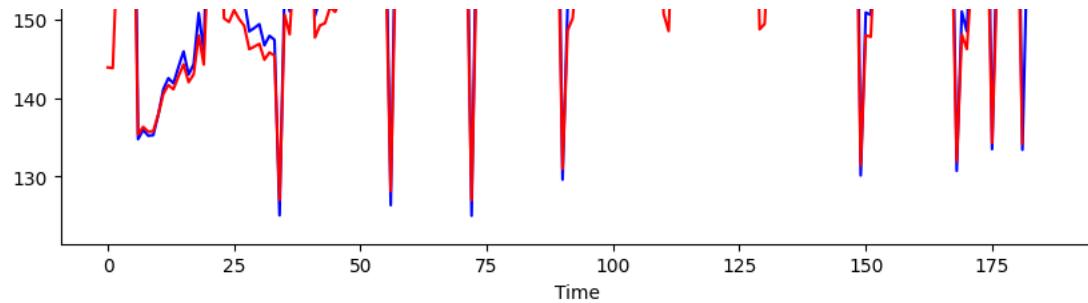
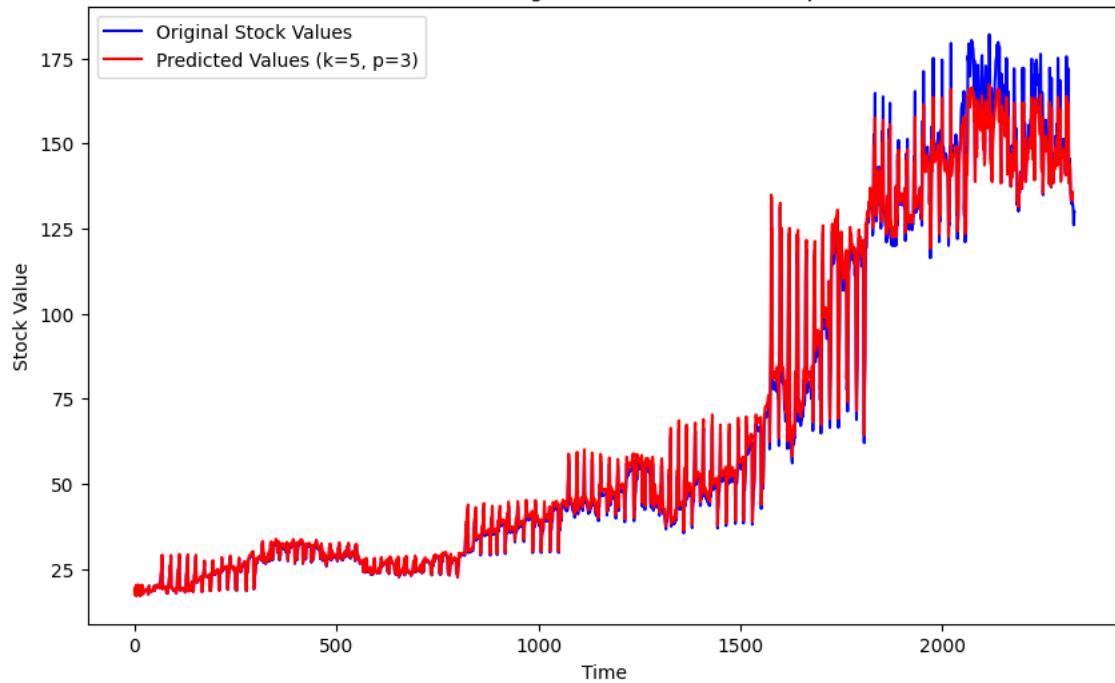
$k = 4, p = 8$
Training error: 121.15508625205123
Testing error: 1903.9779065431433

Predicted vs. Original Stock Values ($k=4, p=8$)Predicted vs. Original Stock Values ($k=4, p=8$)



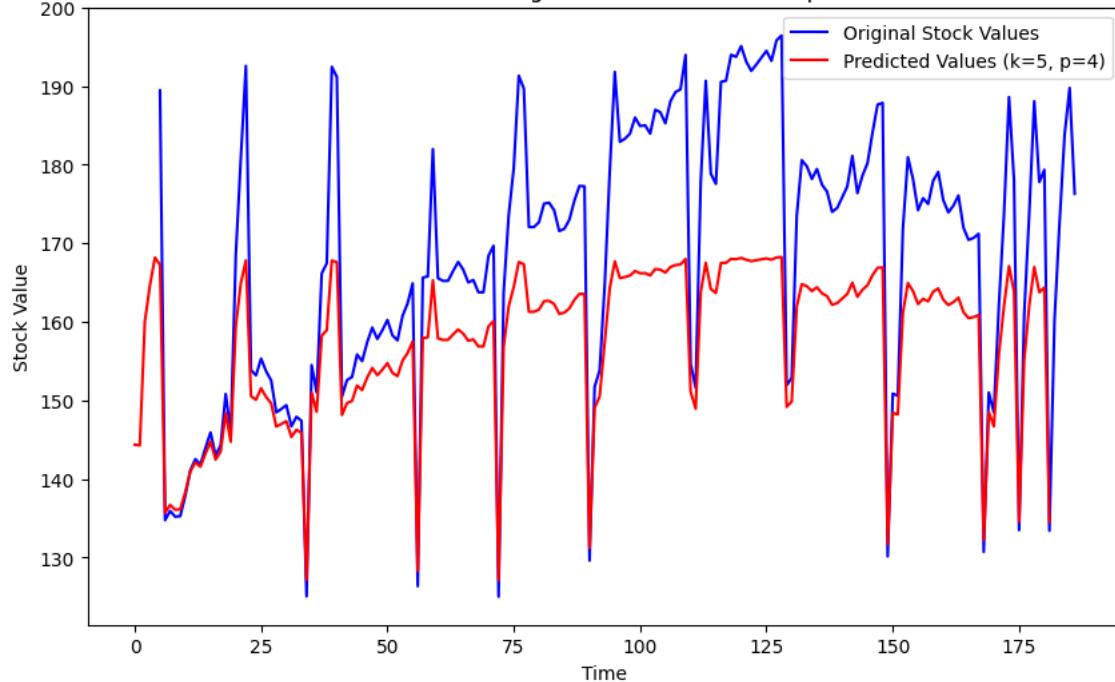
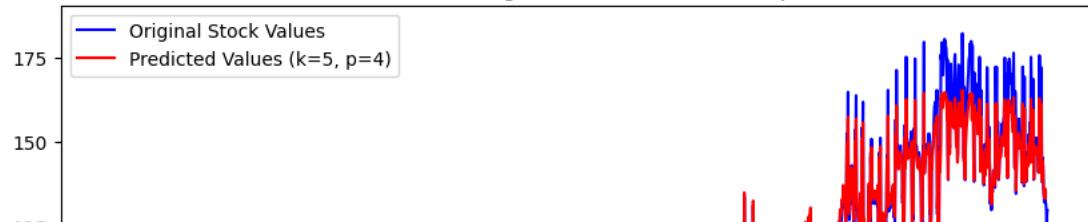
$k = 5, p = 3$
 Training error: 137.05140210193744
 Testing error: 432.42719489079866

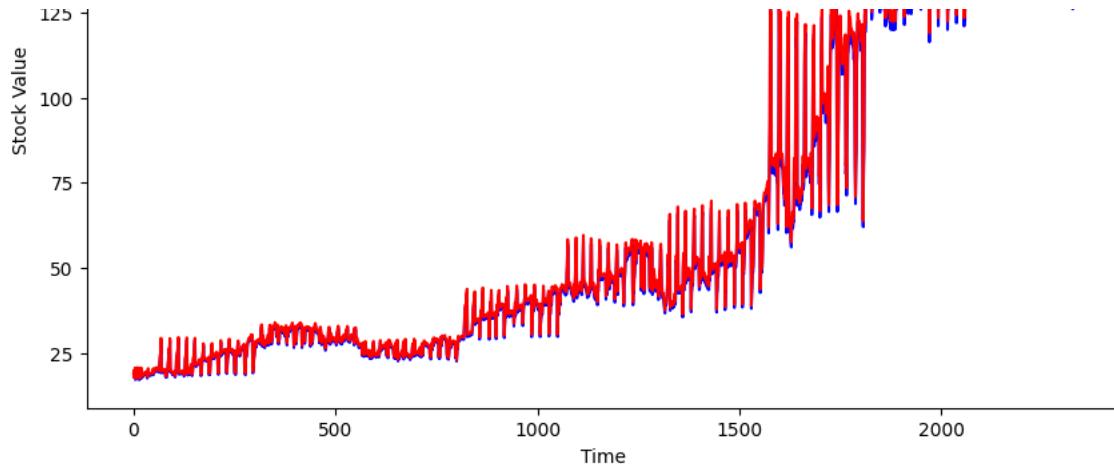


Predicted vs. Original Stock Values ($k=5, p=3$) $k = 5, p = 4$

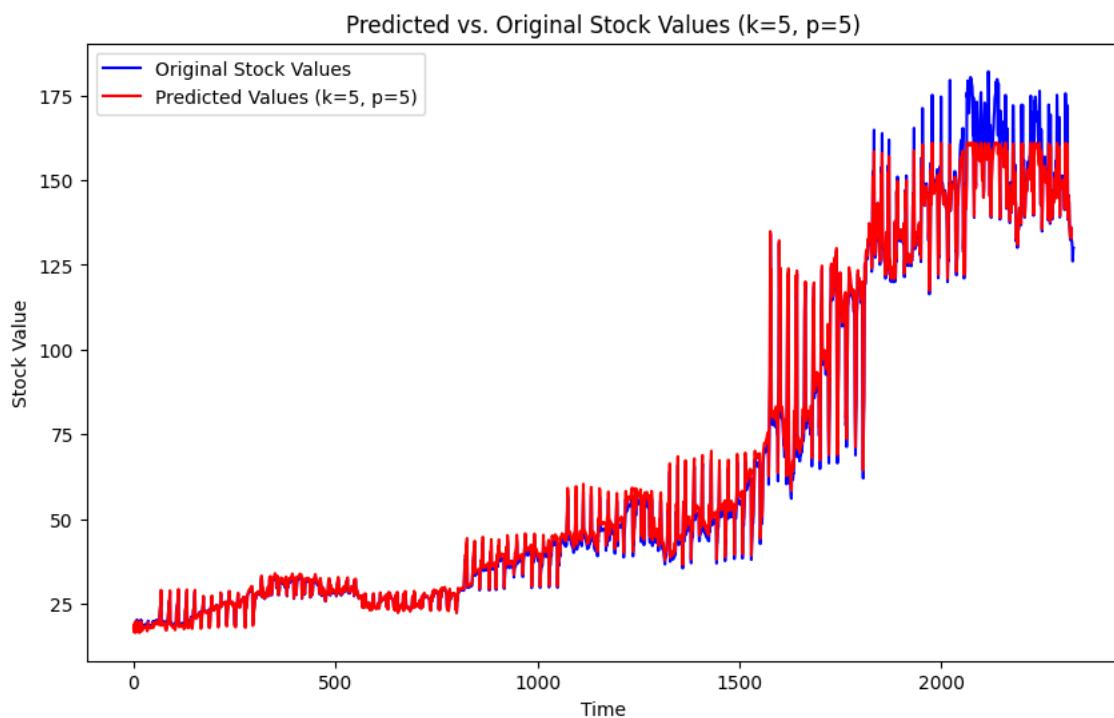
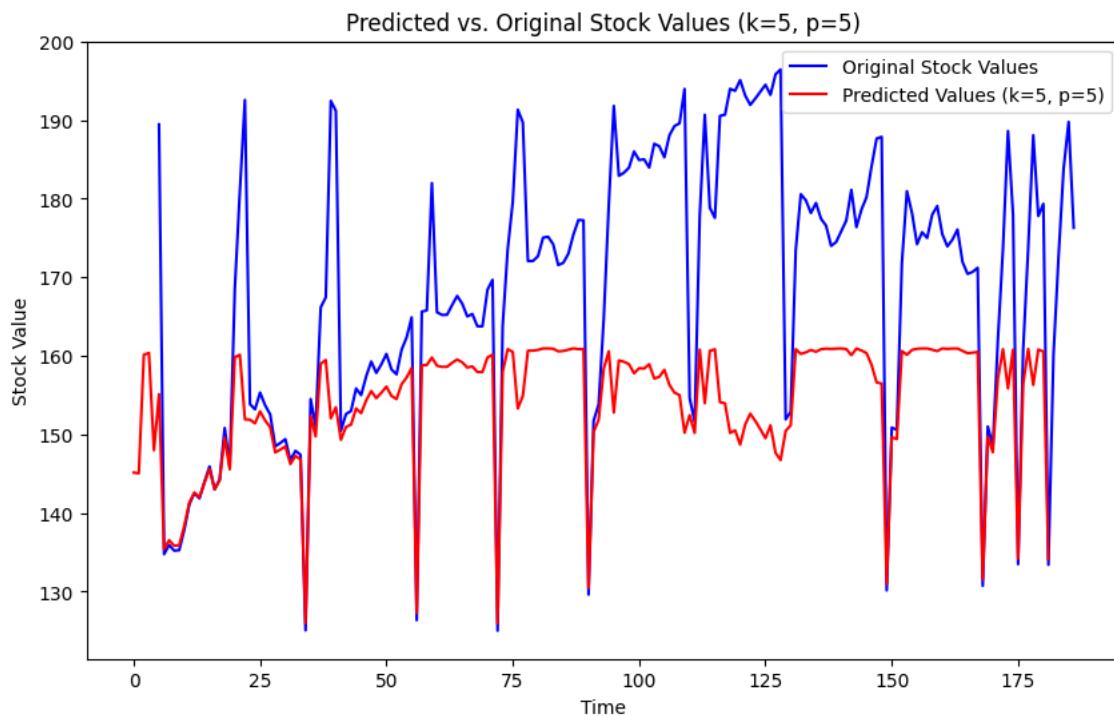
Training error: 136.91004580967368

Testing error: 446.03906288921814

Predicted vs. Original Stock Values ($k=5, p=4$)Predicted vs. Original Stock Values ($k=5, p=4$)

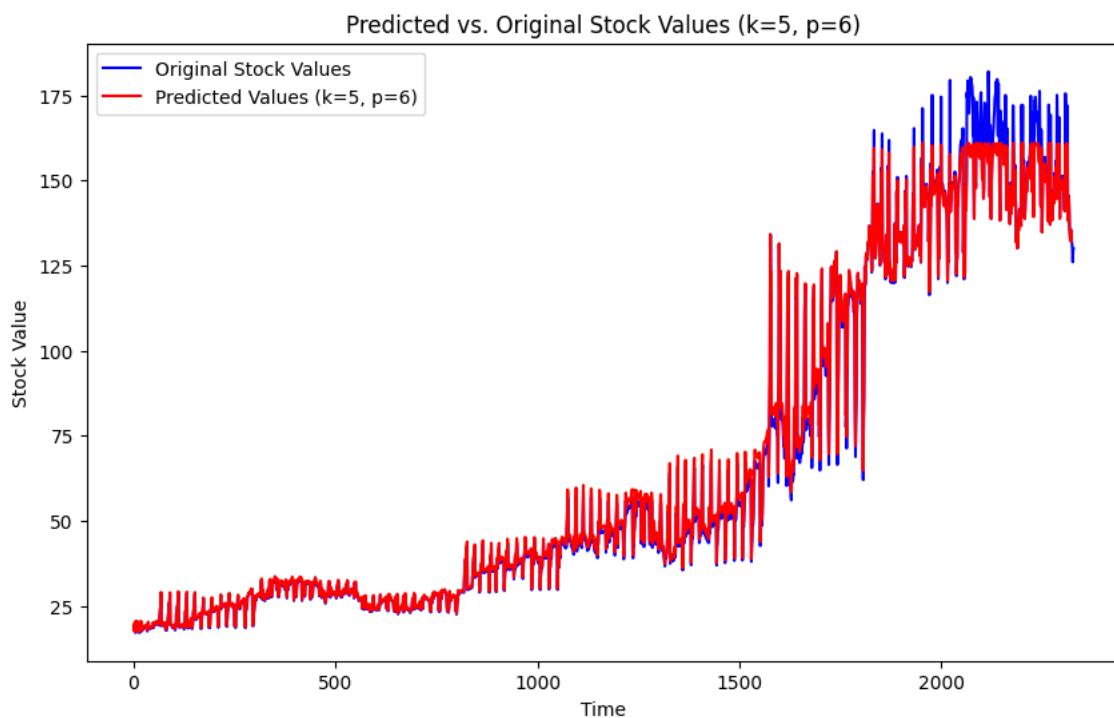
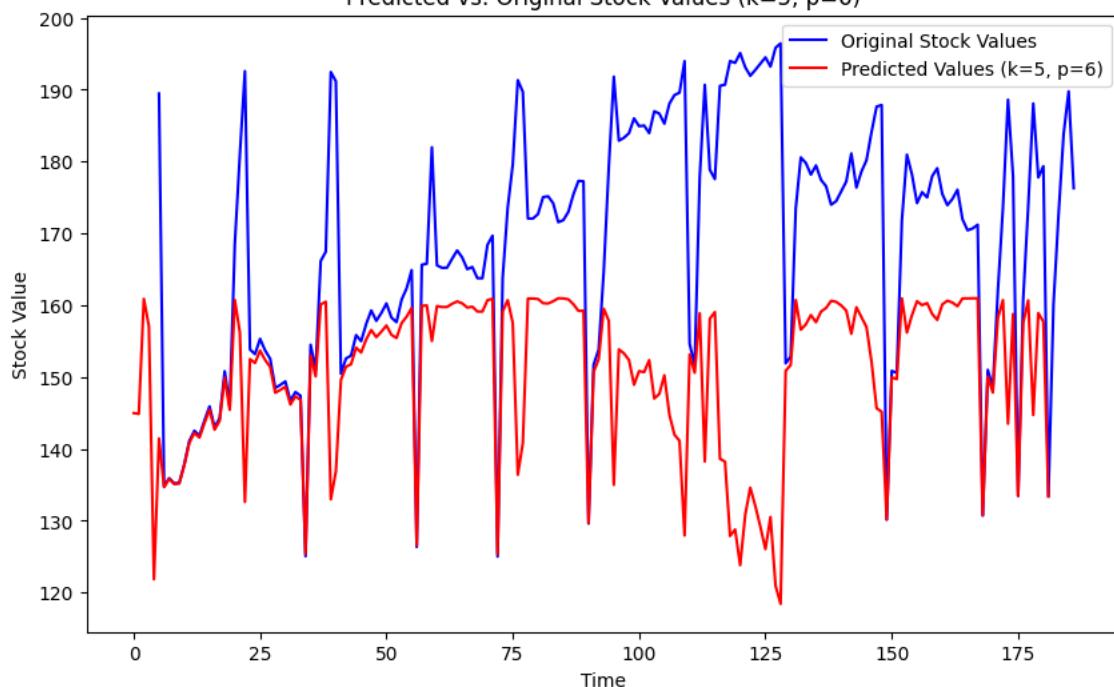


$k = 5, p = 5$
 Training error: 136.2681077957716
 Testing error: 553.1790501934859

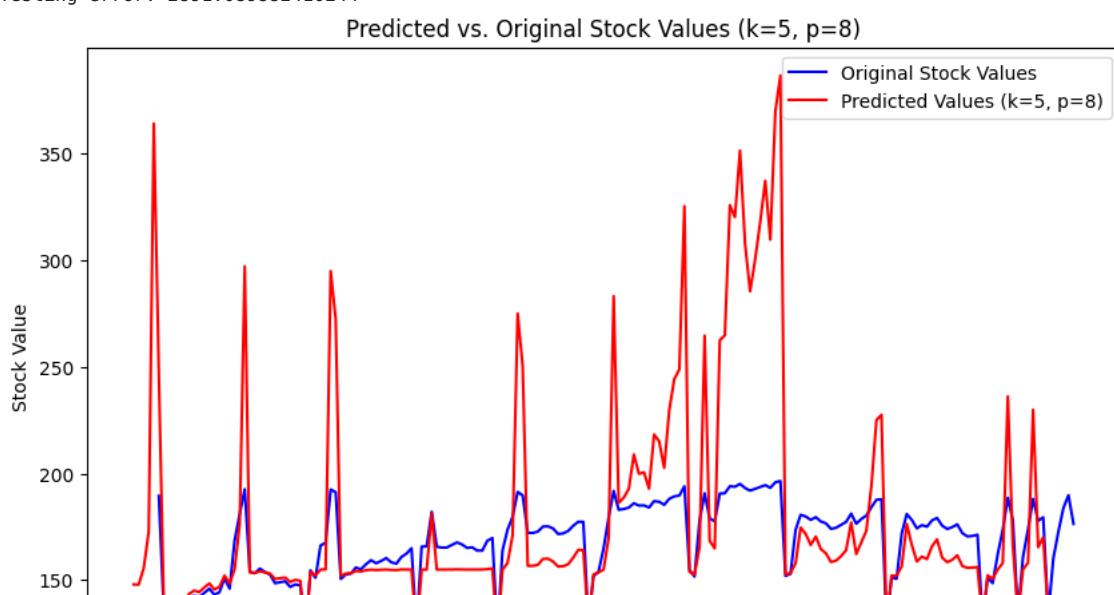


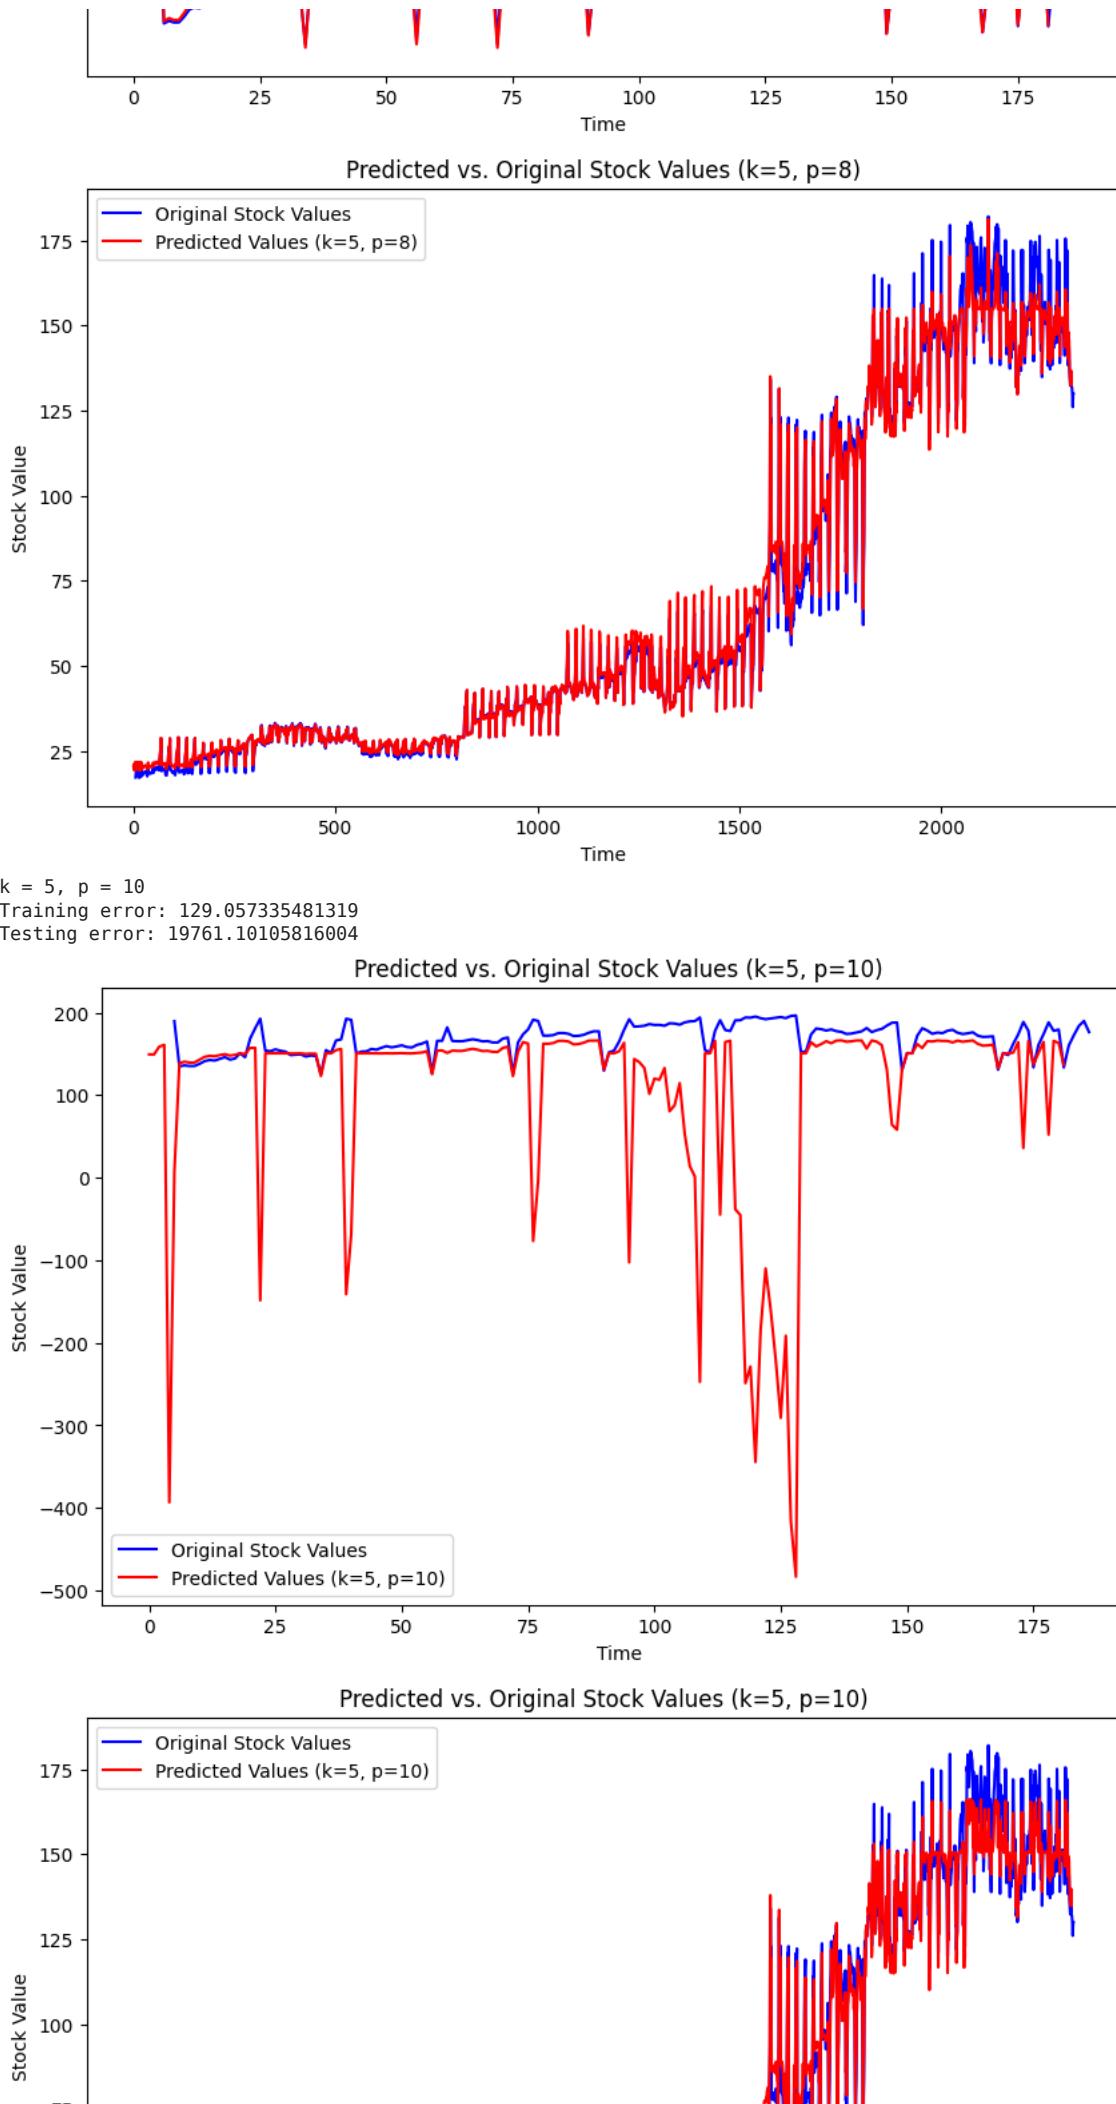
$k = 5, p = 6$
 Training error: 135.92434576648154
 Testing error: 791.3744405195356

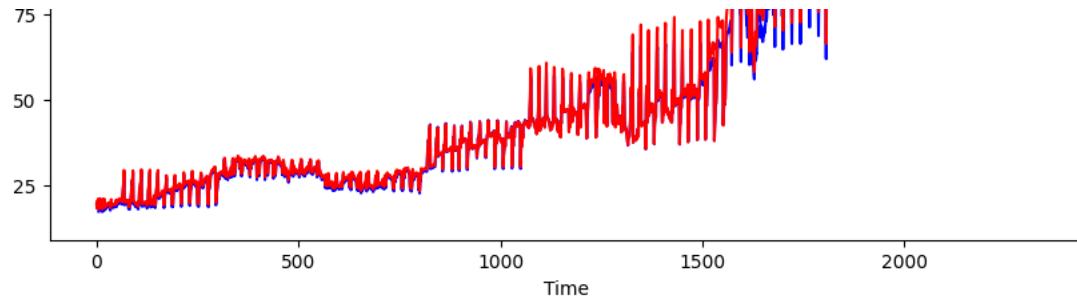
Predicted vs. Original Stock Values ($k=5, p=6$)



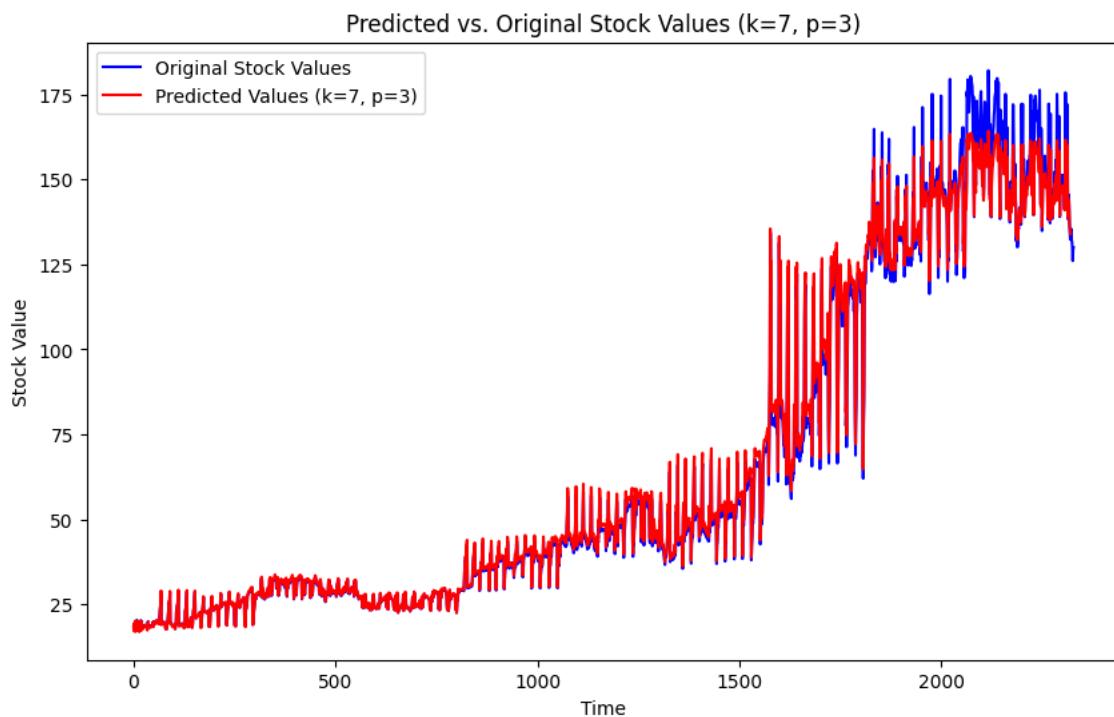
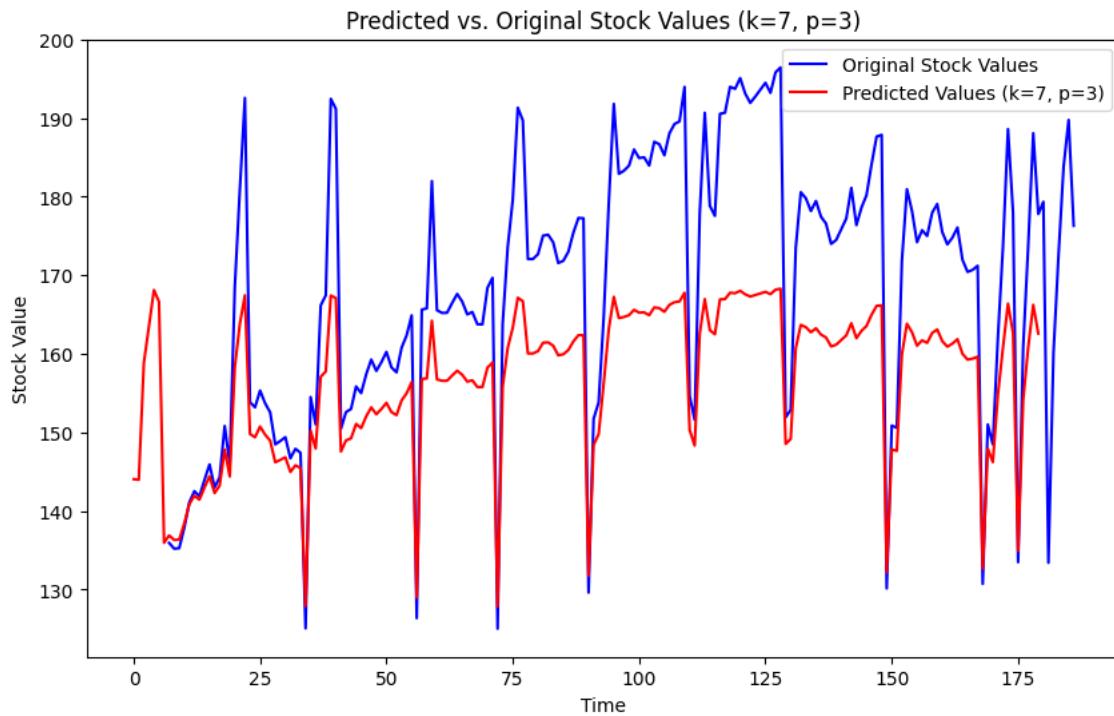
$k = 5, p = 8$
 Training error: 132.23932793646625
 Testing error: 2891.089882410244



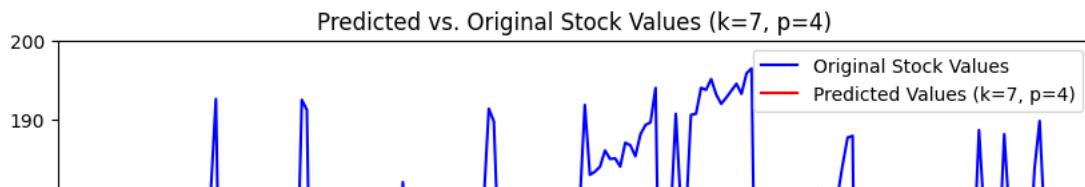


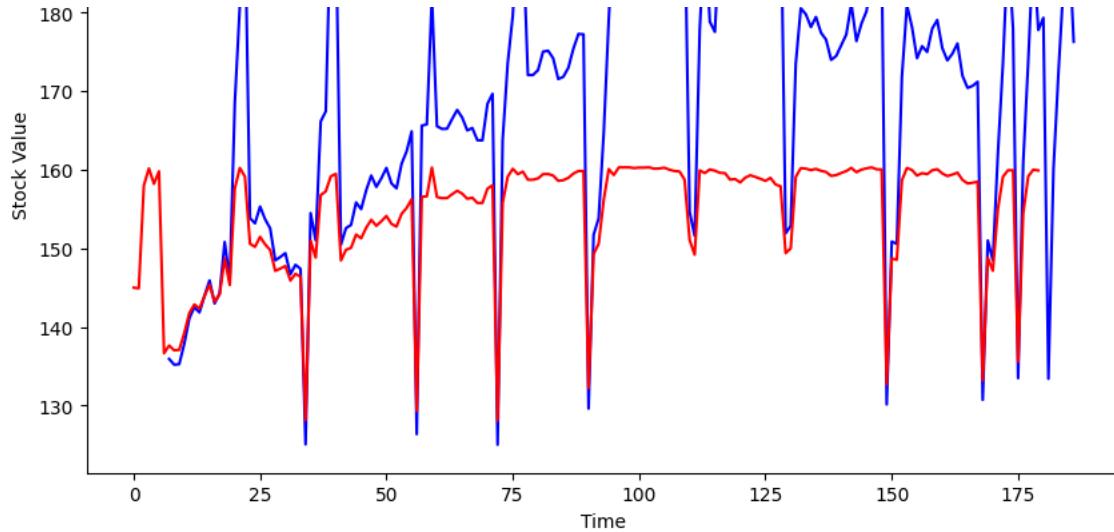
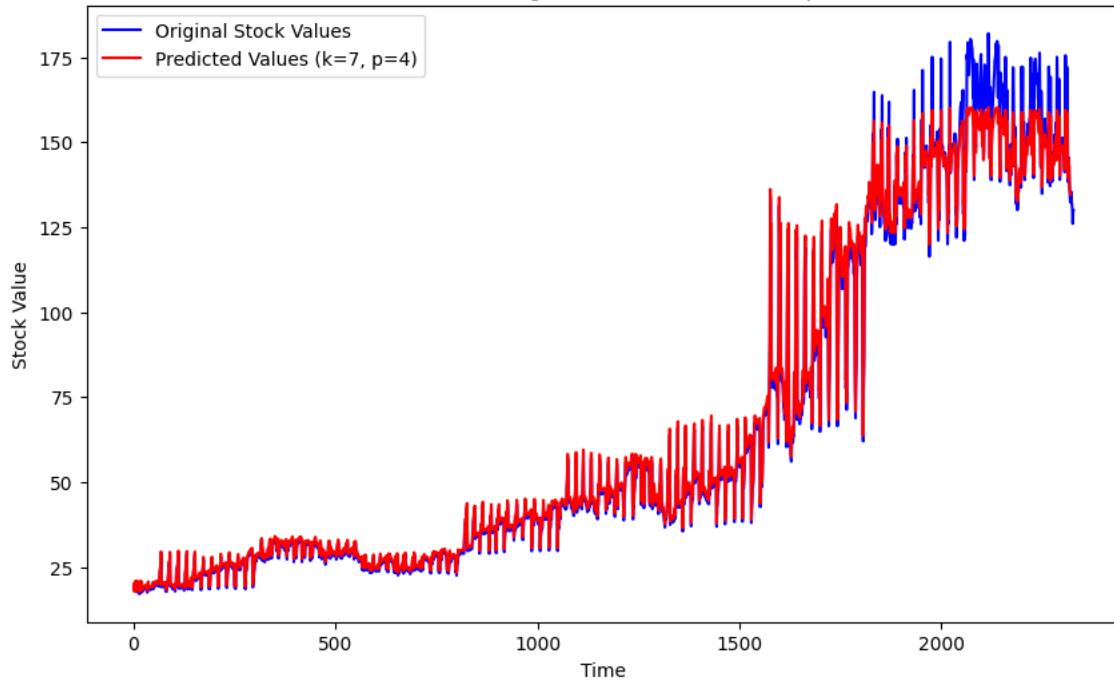


$k = 7, p = 3$
 Training error: 147.08306452605268
 Testing error: 424.16167464666813



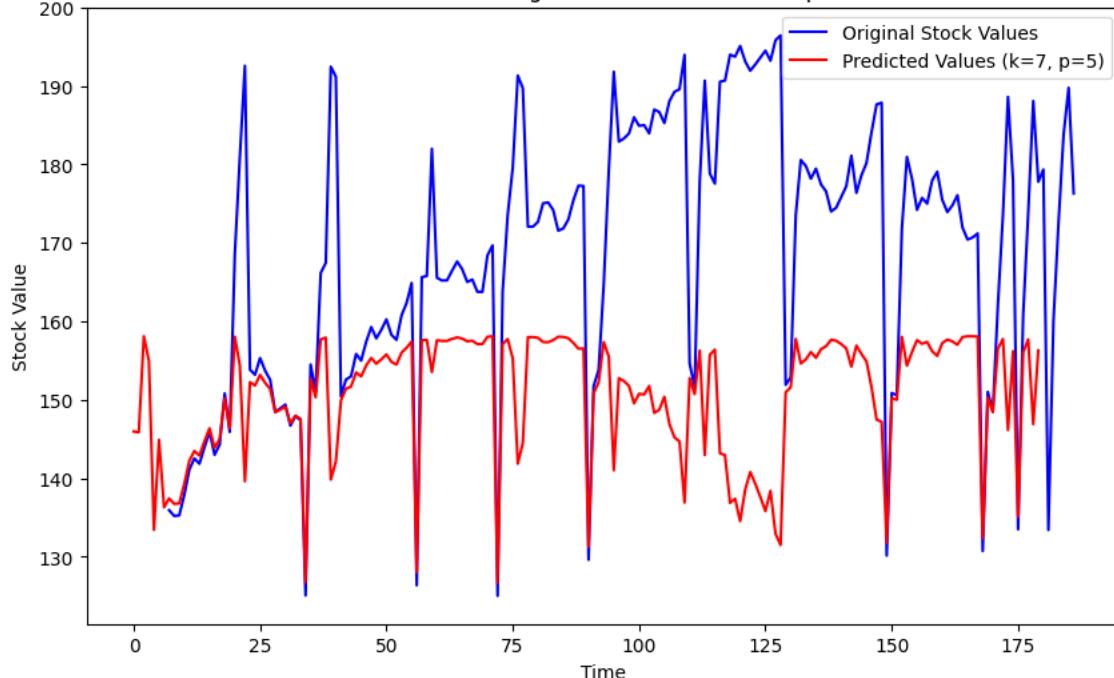
$k = 7, p = 4$
 Training error: 146.51220375981543
 Testing error: 474.3940527130634

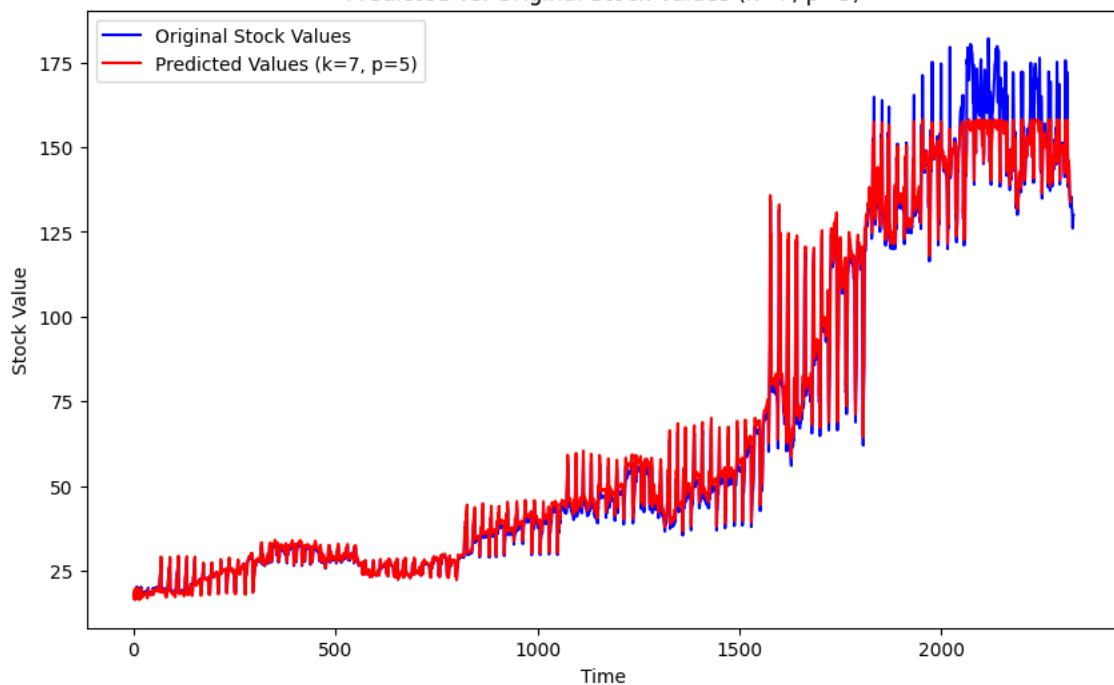


Predicted vs. Original Stock Values ($k=7, p=4$) $k = 7, p = 5$

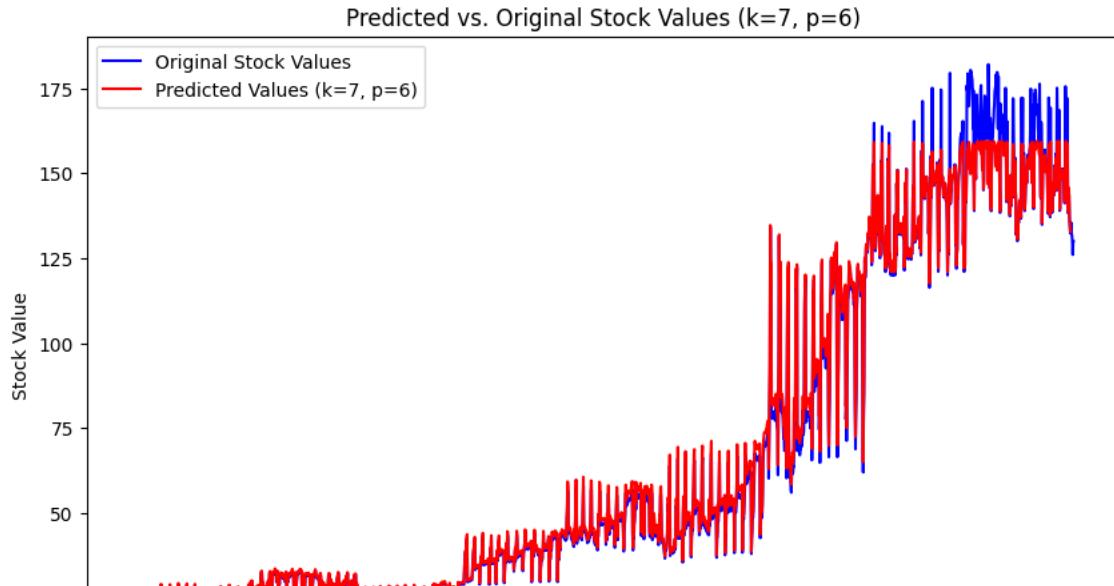
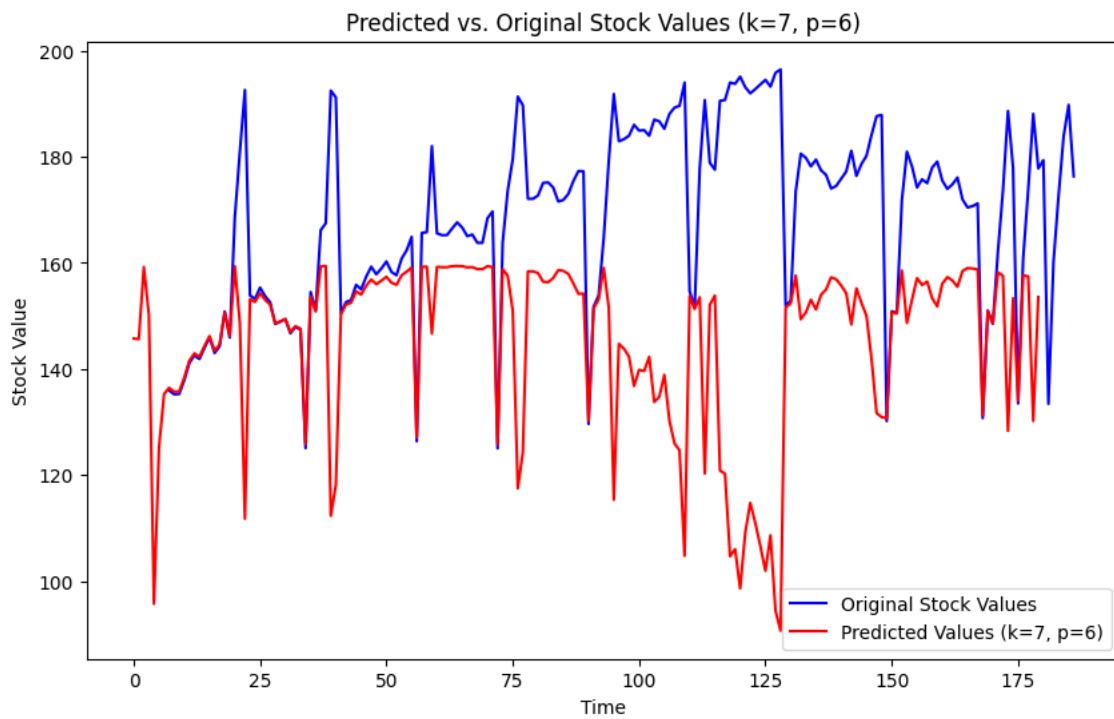
Training error: 145.54301588529881

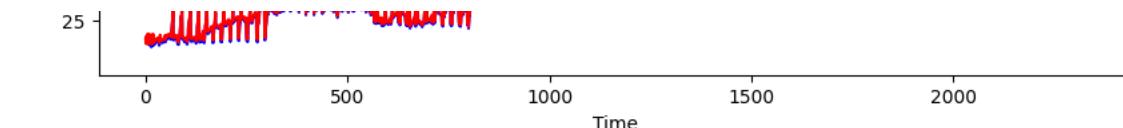
Testing error: 679.1821123362824

Predicted vs. Original Stock Values ($k=7, p=5$)Predicted vs. Original Stock Values ($k=7, p=5$)

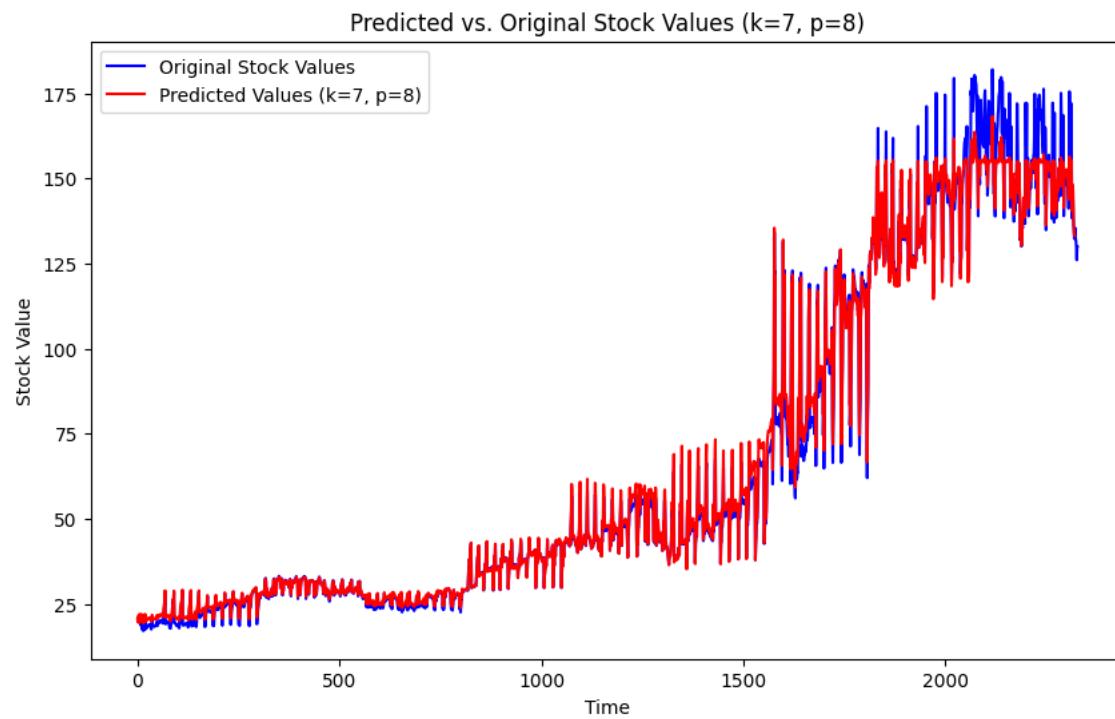
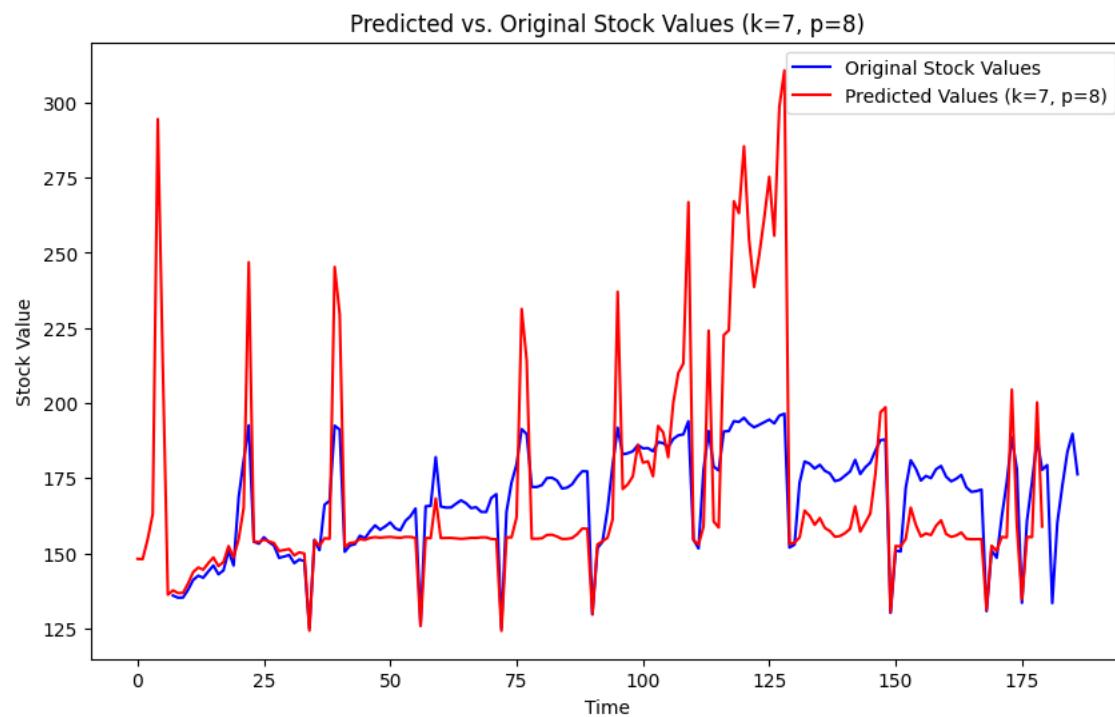


$k = 7, p = 6$
Training error: 144.82661952889802
Testing error: 1185.0629165045127

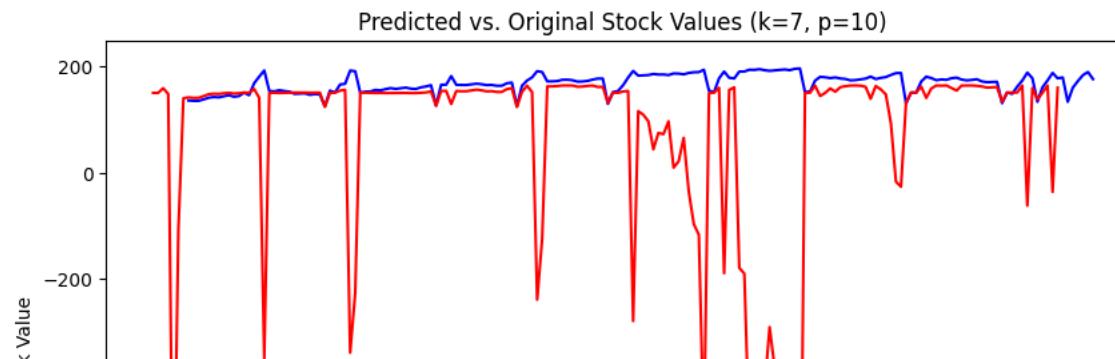


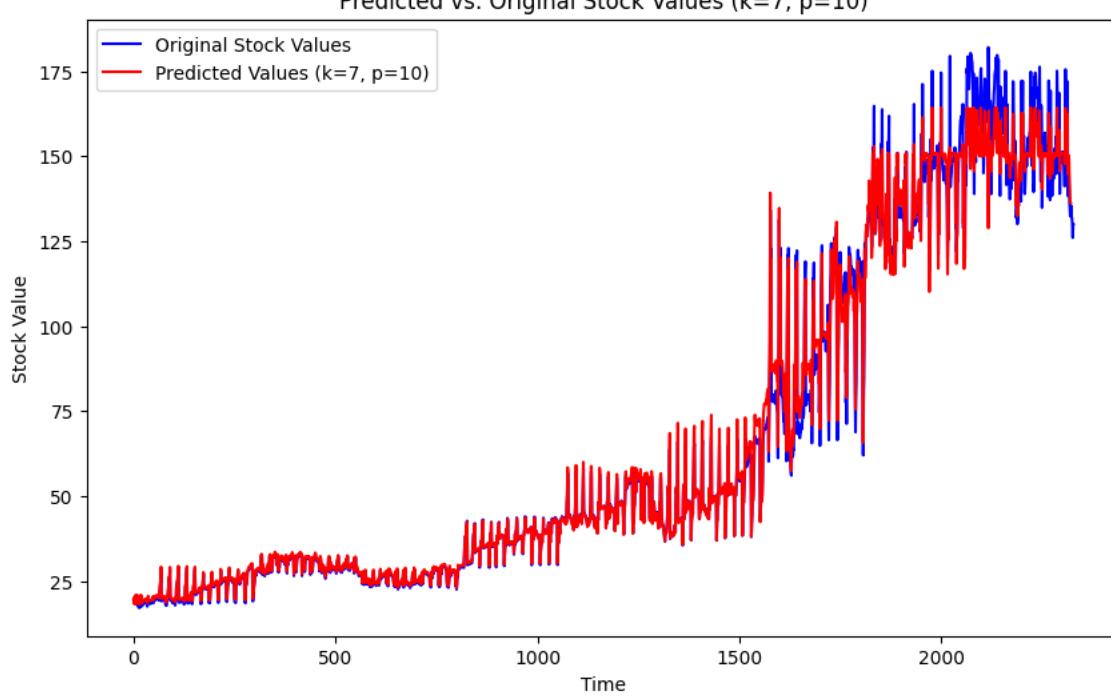
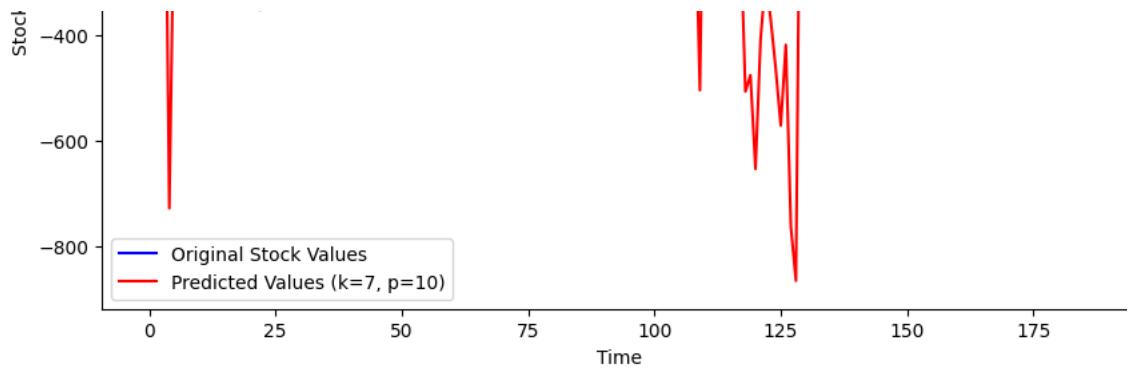


$k = 7, p = 8$
 Training error: 142.28493841076704
 Testing error: 1282.4525167145798

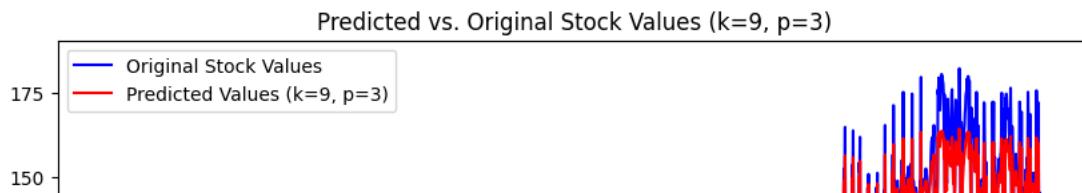
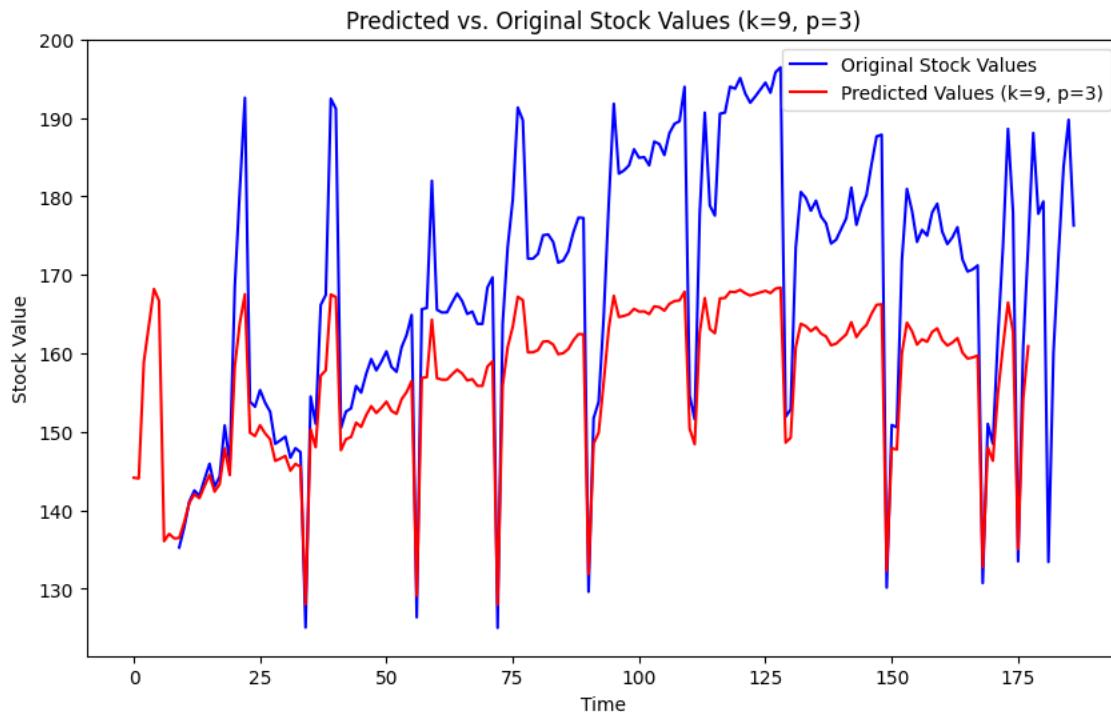


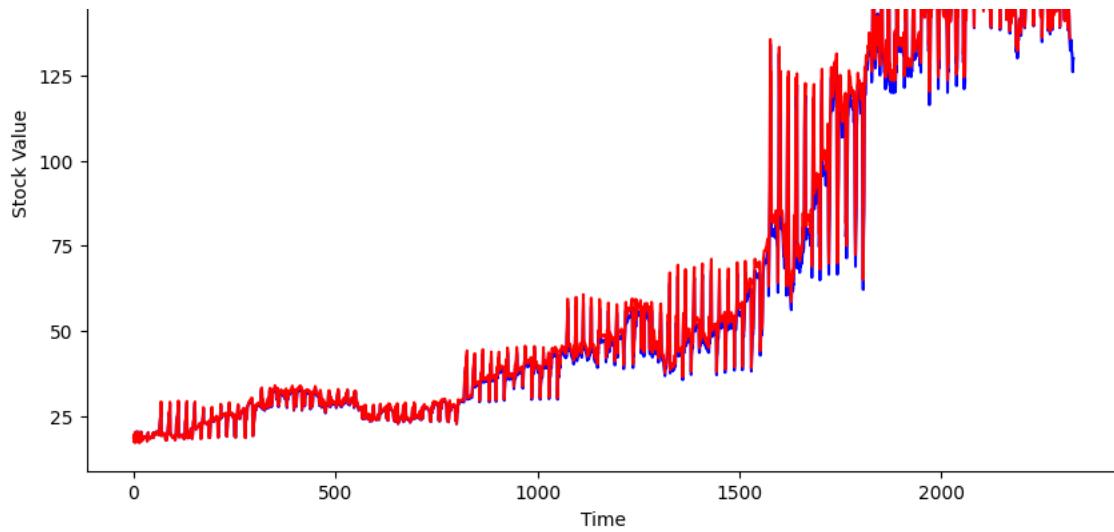
$k = 7, p = 10$
 Training error: 137.0845242016022
 Testing error: 51447.62352476746



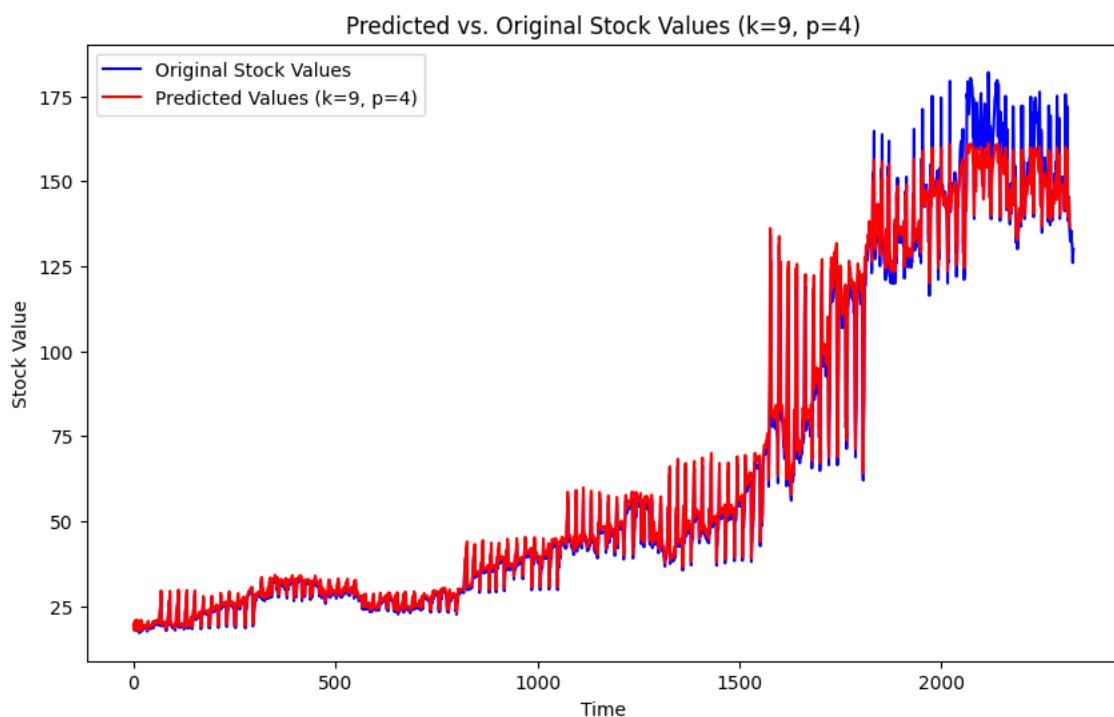
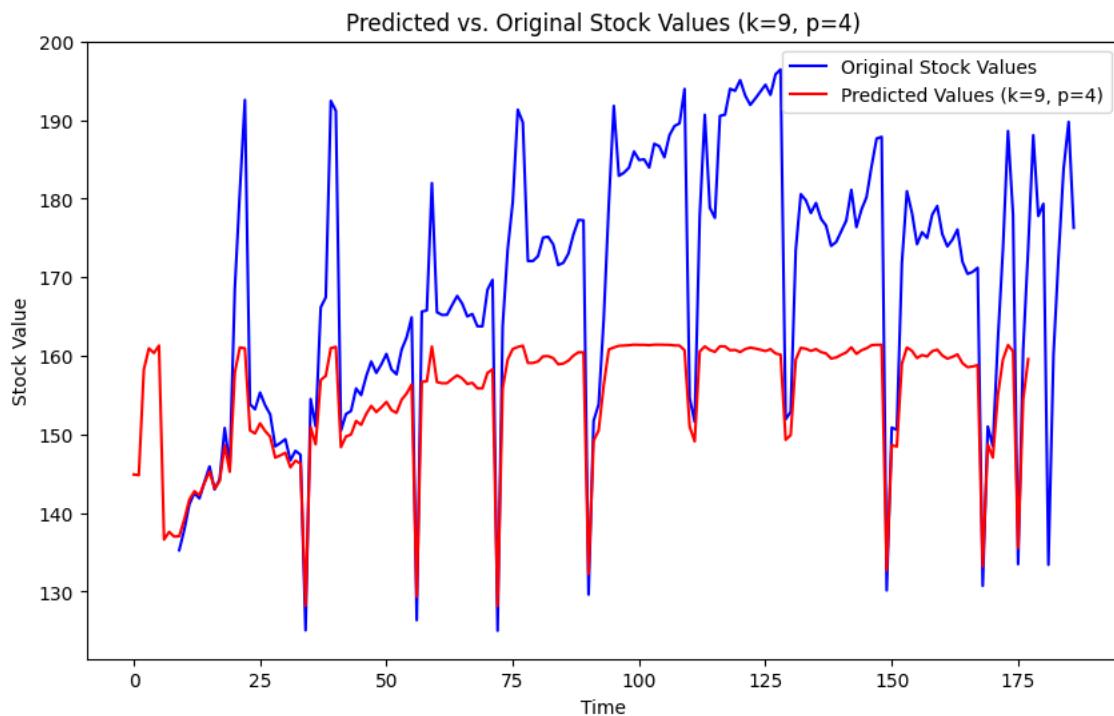


$k = 9, p = 3$
Training error: 147.51041657593643
Testing error: 462.86853007685636



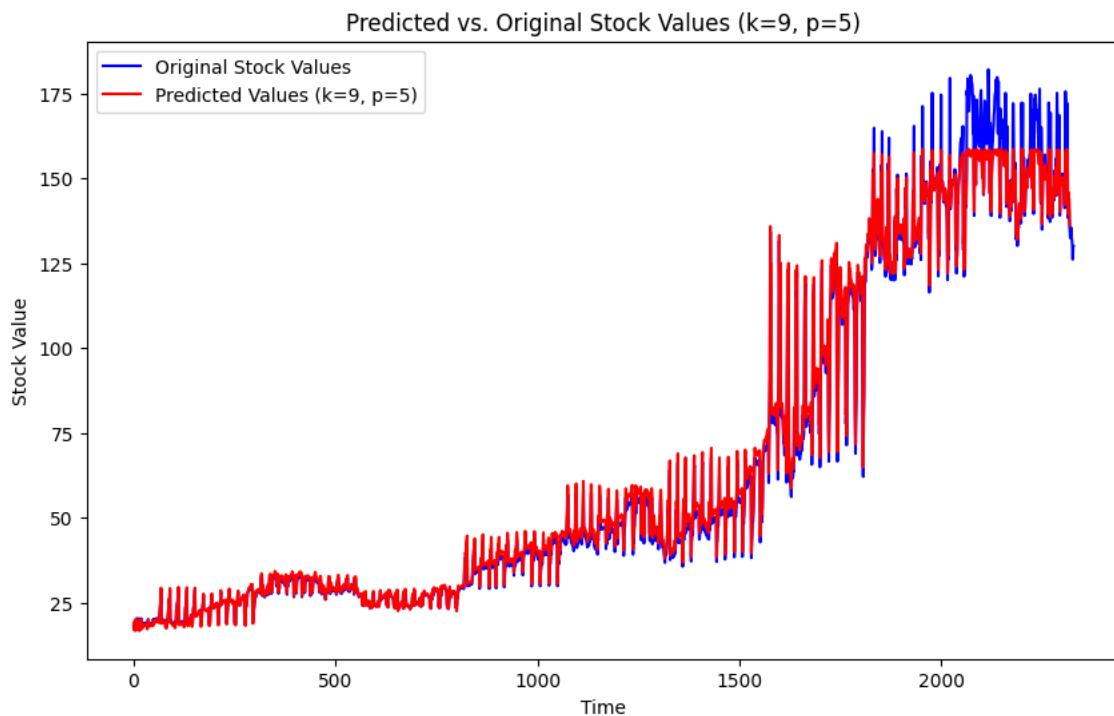
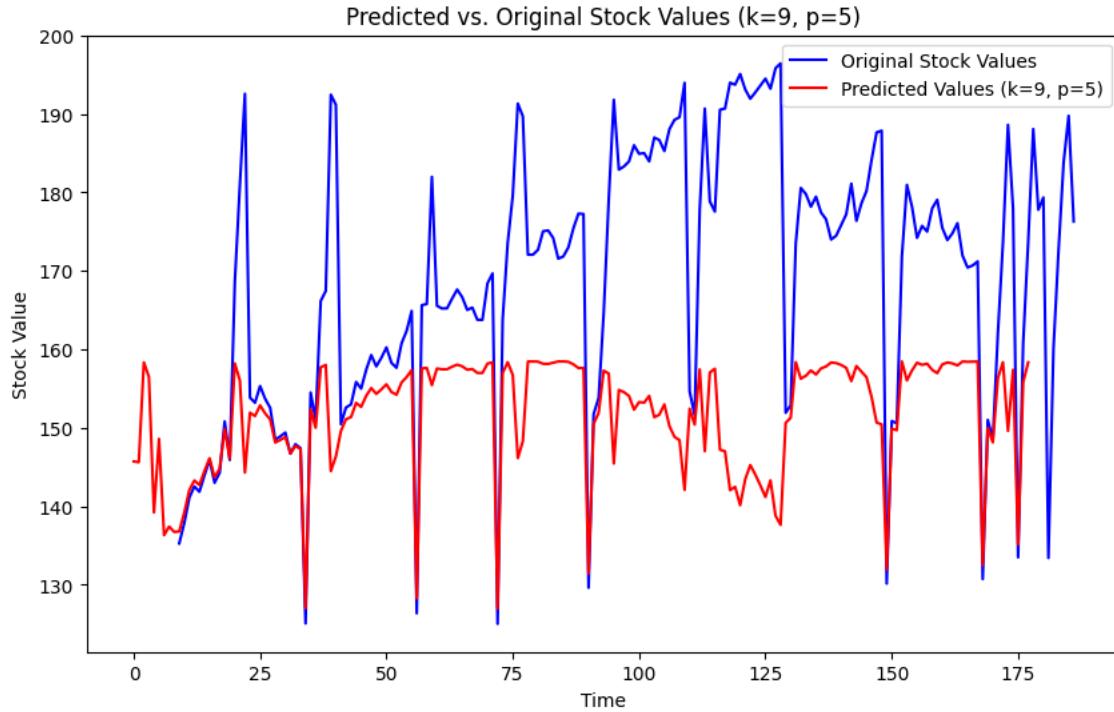


$k = 9, p = 4$
 Training error: 147.15224721941155
 Testing error: 494.31711262227753

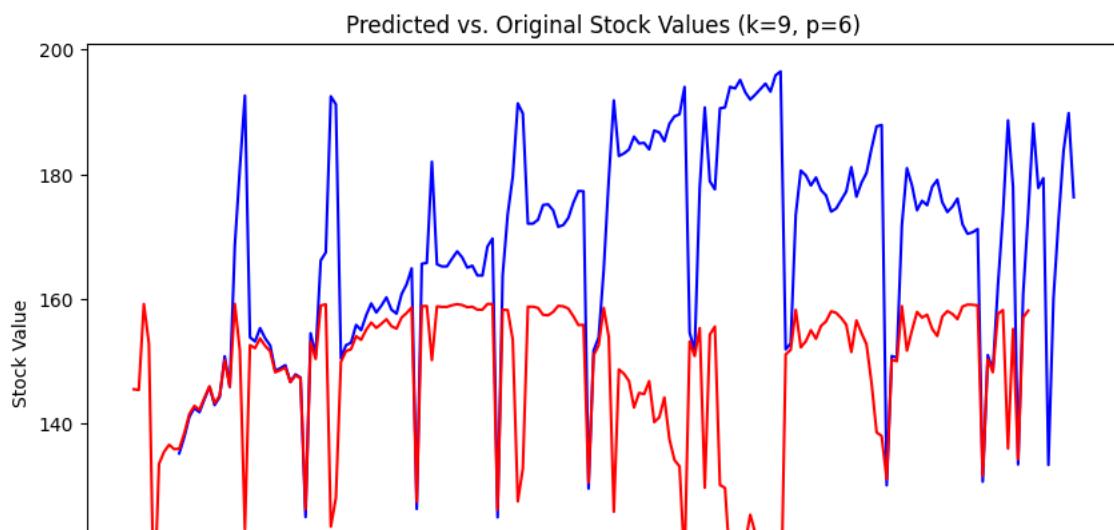


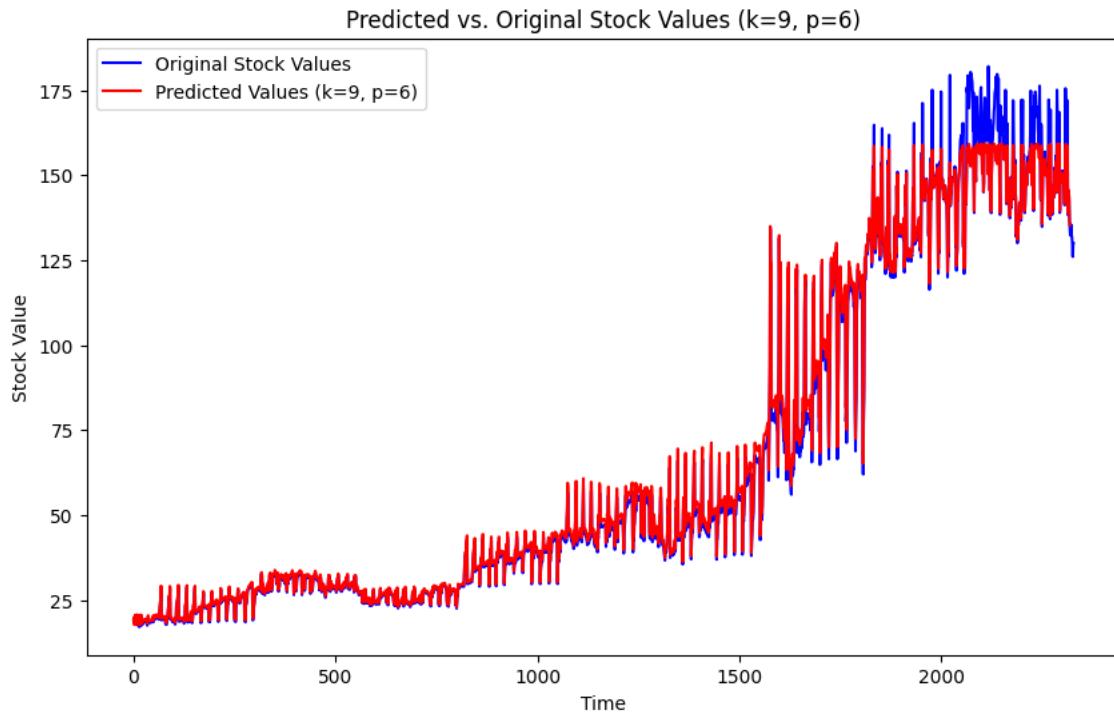
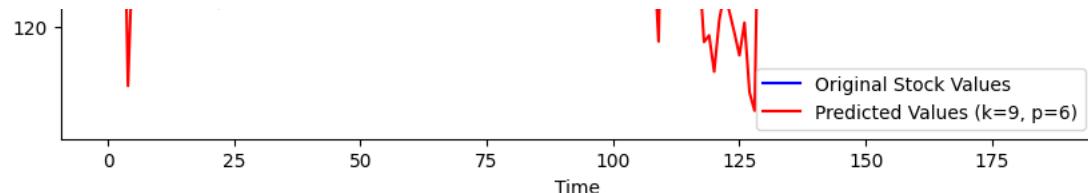
$k = 9, p = 5$
 Training error: 146.1165157179001

Training error: 140.440545 / 170054
 Testing error: 637.2384949844766



$k = 9, p = 6$
 Training error: 146.0315229583377
 Testing error: 949.9641165345181





$k = 9, p = 8$
 Training error: 143.06111176411676
 Testing error: 1764.893181985691

